

CS 4602

Introduction to Machine Learning

Advanced Deep Networks and
Learning Strategies

Instructor: Po-Chih Kuo

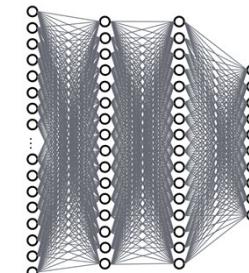
Roadmap

- Introduction and Basic Concepts
- Regression
- Bayesian Classifiers
- Decision Trees
- KNN
- Linear Classifier
- Neural Networks
- Deep learning
- Convolutional Neural Networks
- RNN/Transformer
- Reinforcement Learning
- Model Selection and Evaluation
- Clustering
- Data Exploration & Dimensionality reduction

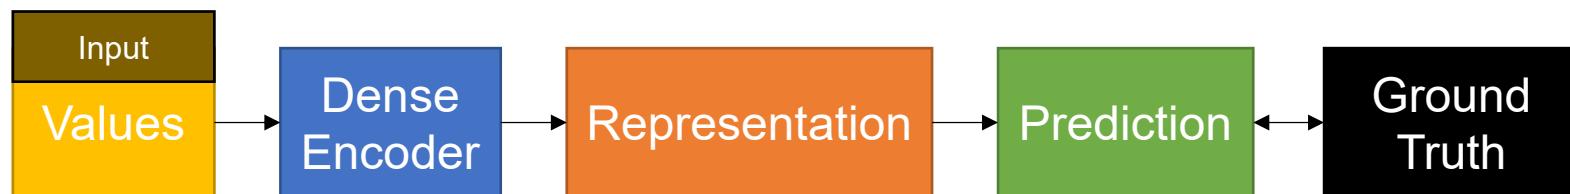
Learning, learning, learning

- Supervised learning
 - FFNNs
 - CNNs
 - RNNs
 - Encoder Decoder
 - Unsupervised learning
 - Autoencoder
 - GAN
 - Diffusion model
 - Self-Supervised learning
 - Large language model
 - Reinforcement learning
 - Transfer learning
 - Federated learning
- 
- Generative AI (GAI)

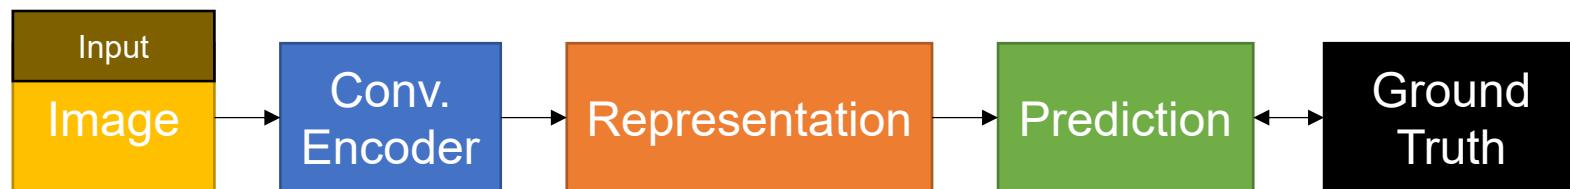
Supervised learning



- **Feed Forward Neural Networks (FFNNs)** - classification and regression based on features.

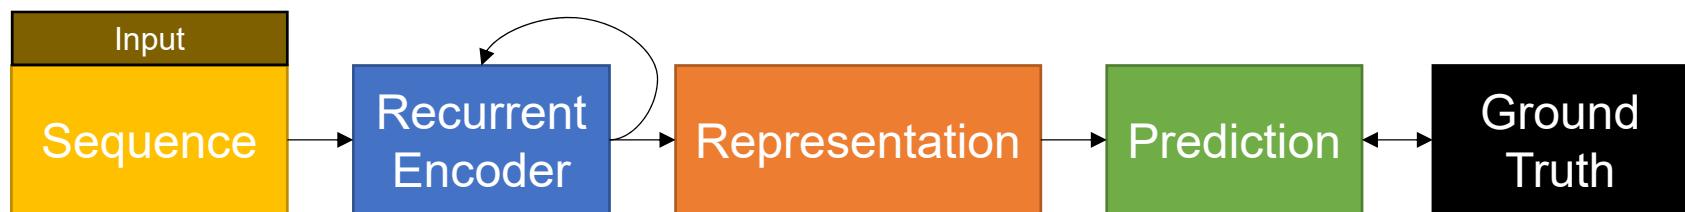


- **Convolutional Neural Networks (CNNs)** - image classification, object detection, video action recognition



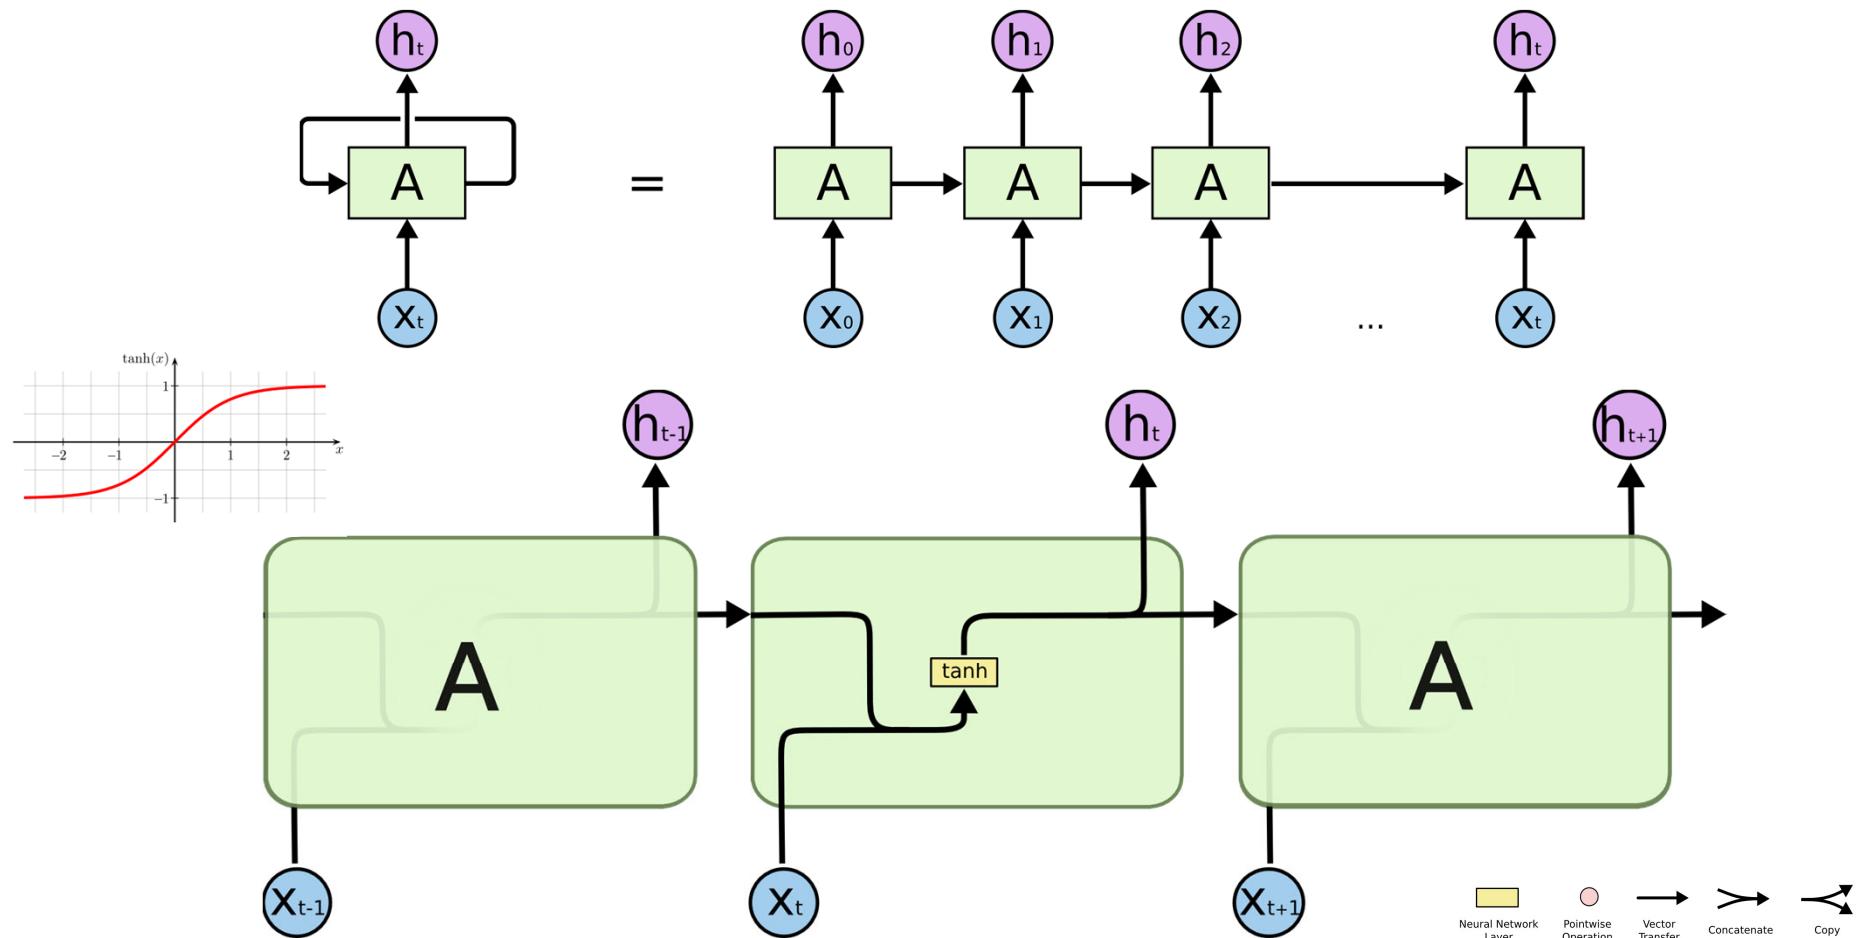
Supervised learning (cont.)

- **Recurrent Neural Networks (RNNs)** - language modeling, speech recognition/generation



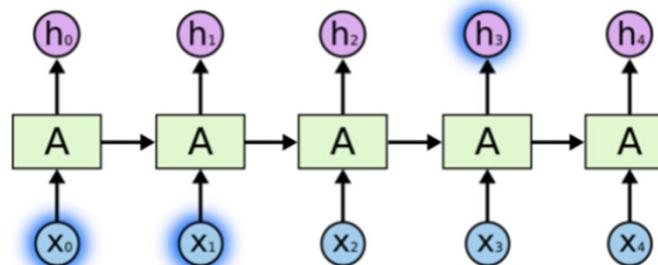
- **Encoder Decoder Architectures** - semantic segmentation, machine translation

Recurrent Neural Network (RNN)

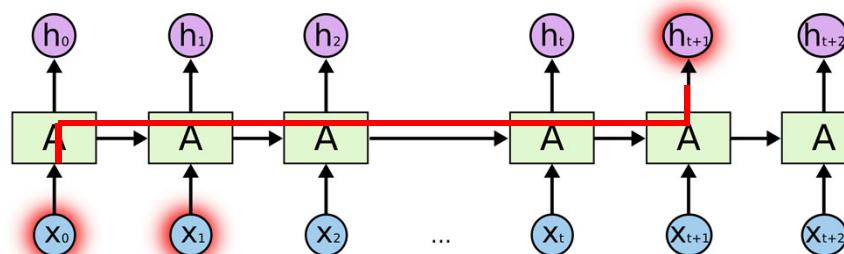


Long-Short Term Dependencies

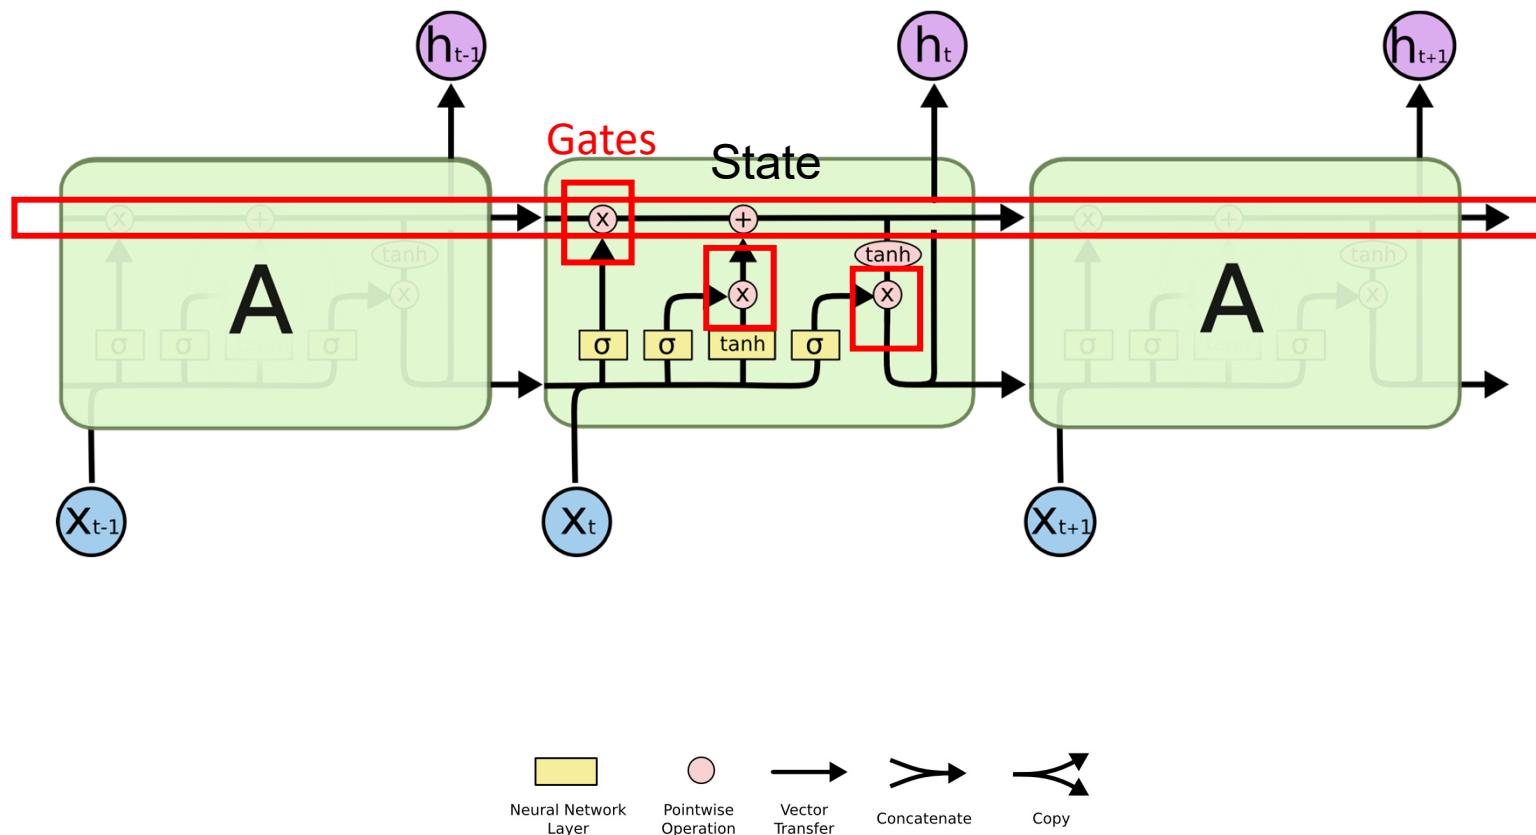
- Short-Term

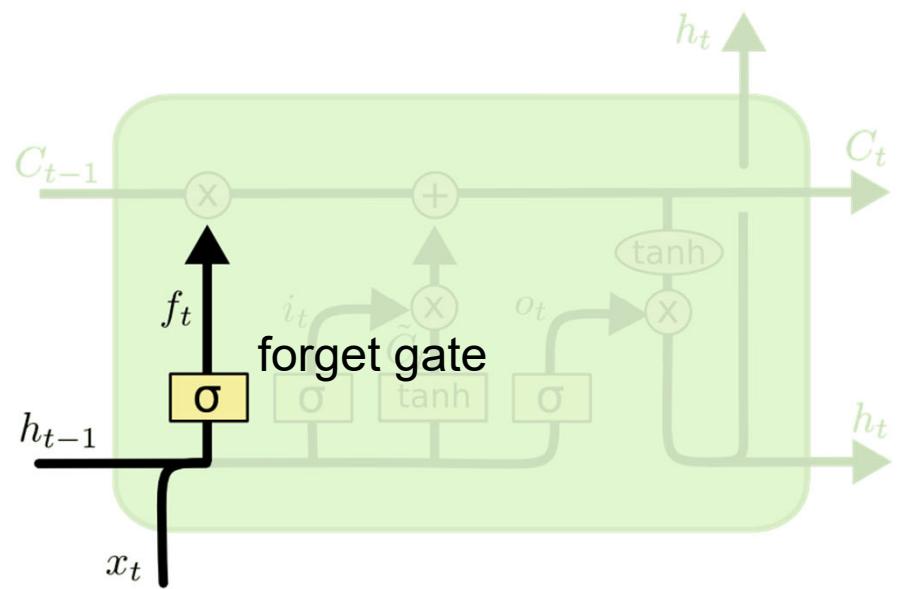


- Long-Term



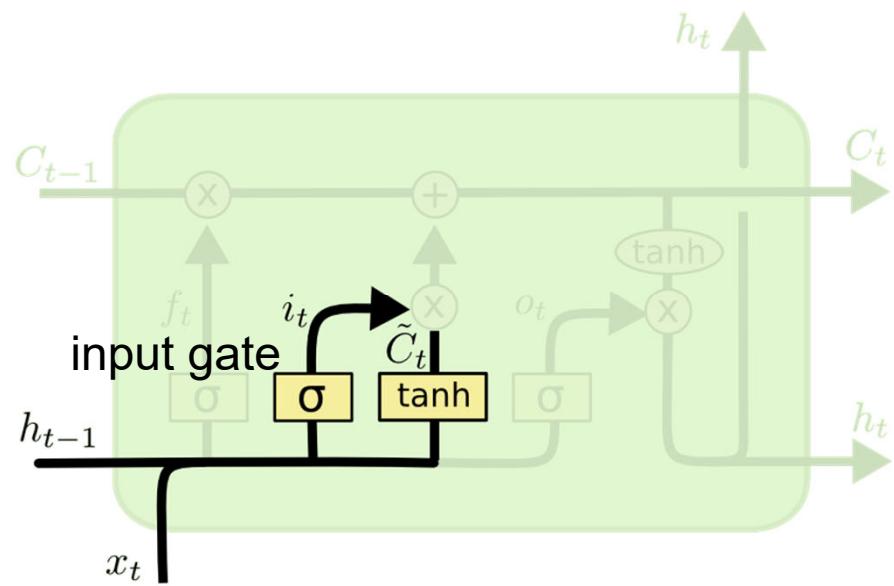
Long-Short Term Memory (LSTM)





$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

forget gate:
1: keep it
0: forget it

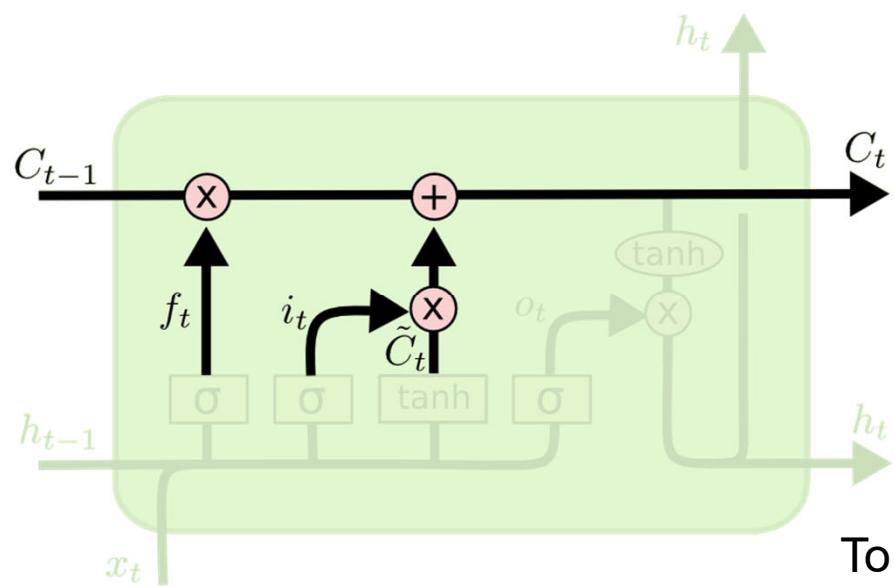


$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

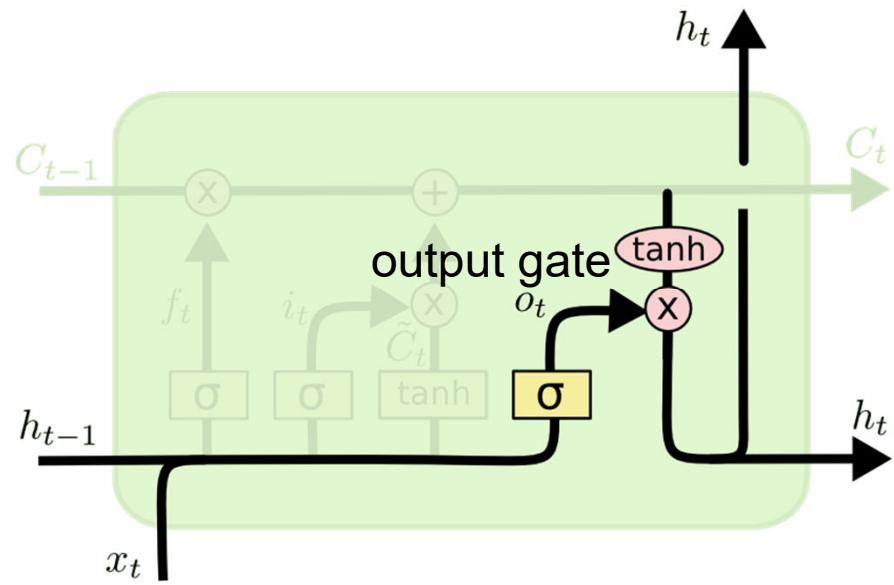
input gate: To decide what new information to store in the cell state.

\tilde{C}_t : new candidate values that could be added to the state.



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

To update the old cell state, C_{t-1} , into the new cell state C_t

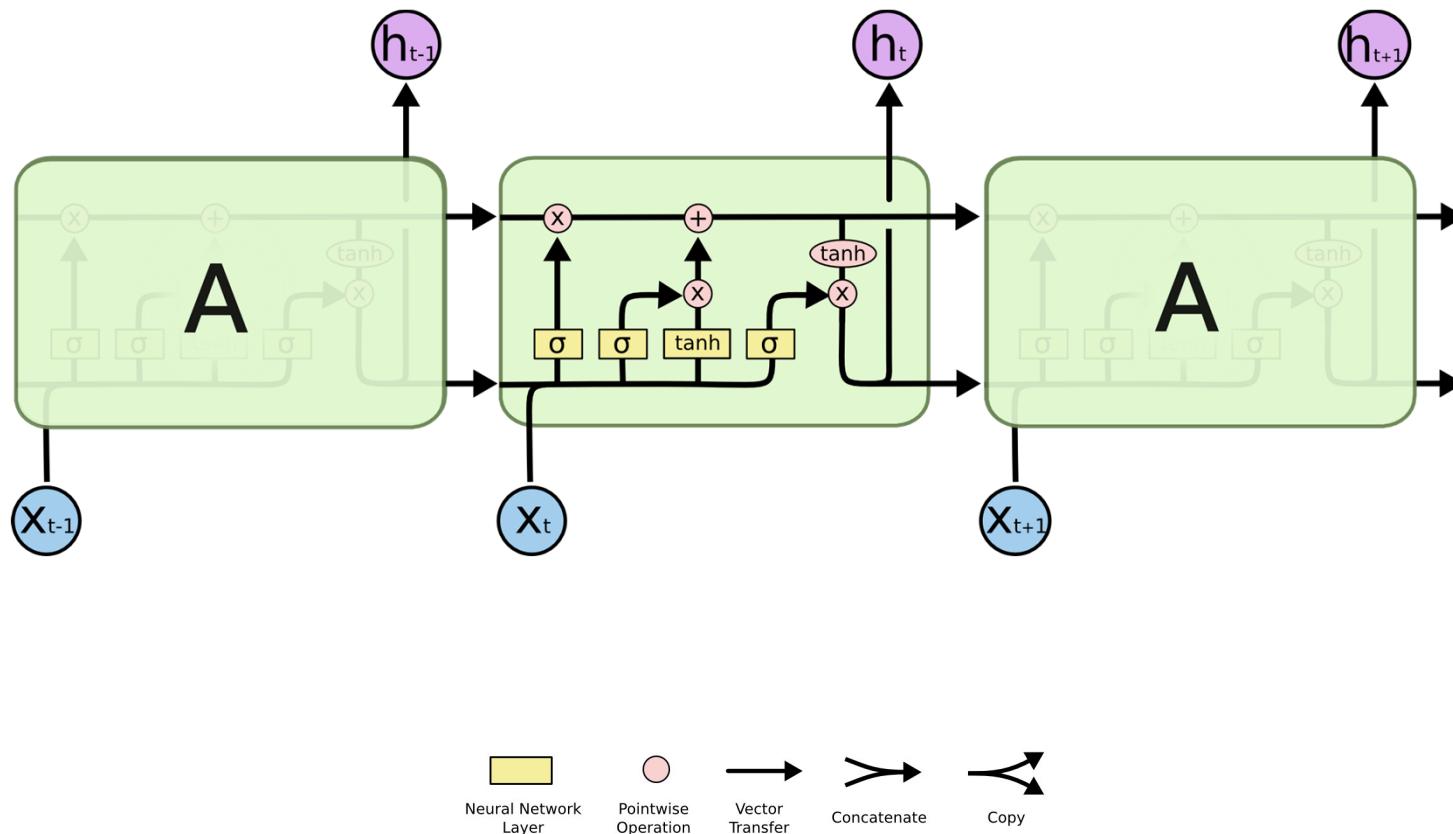


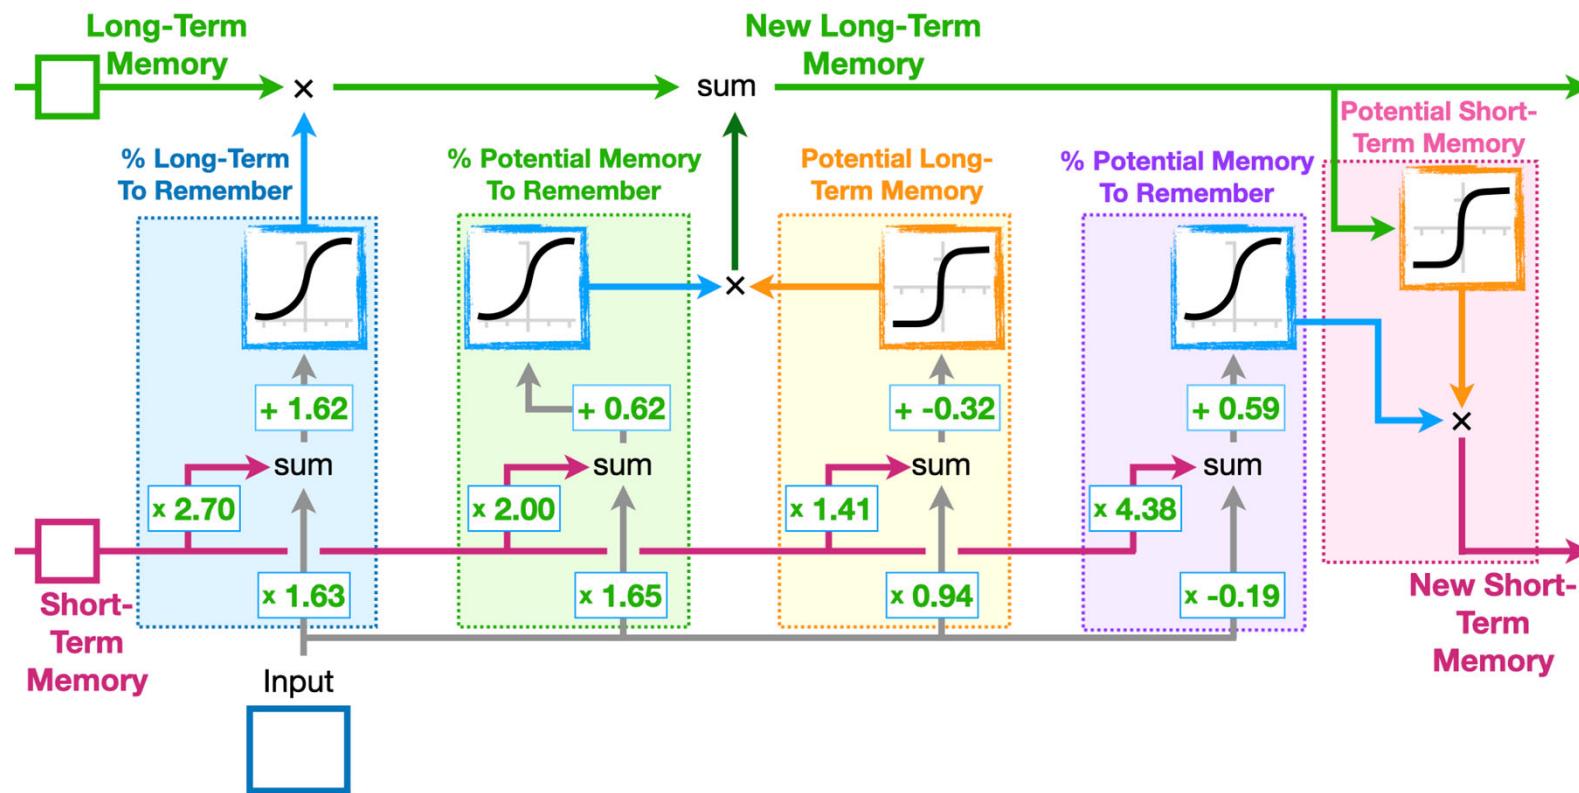
$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$

To decide the output

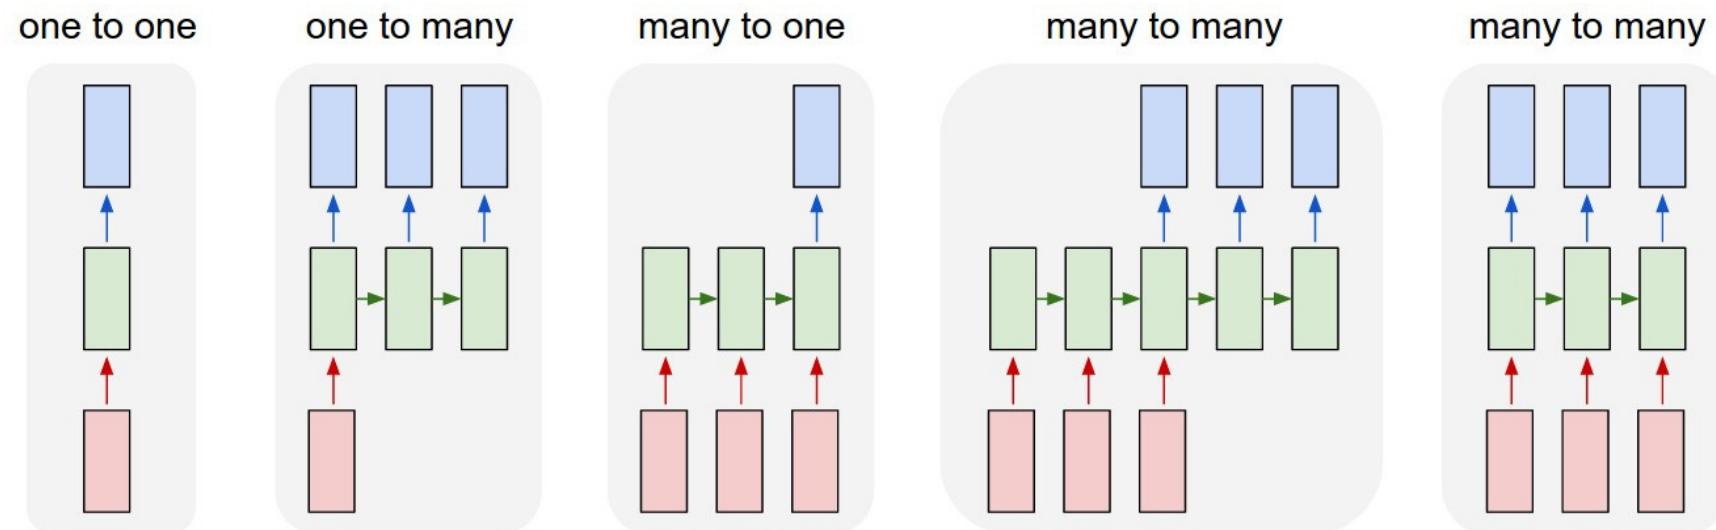
LSTM





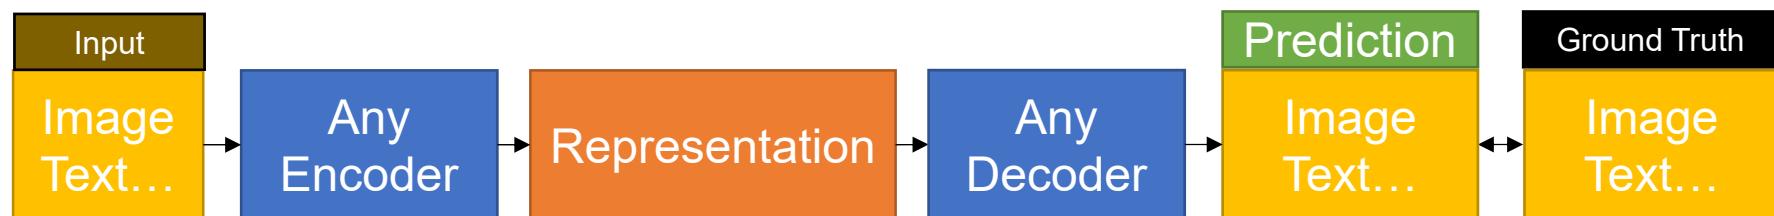
Ref: <https://lightning.ai/lightning-ai/studios/statquest-long-short-term-memory-lstm-with-pytorch-lightning>

Multi-layer LSTM



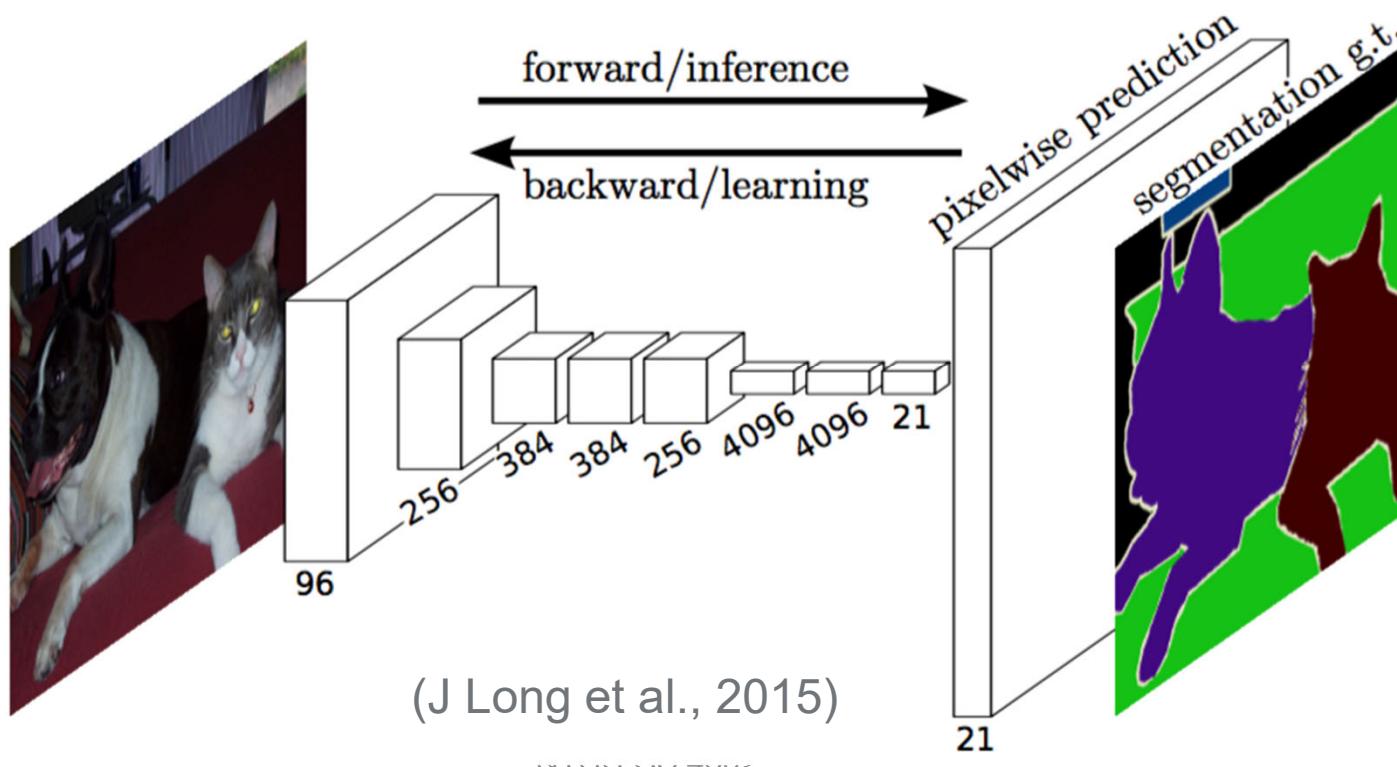
Supervised learning (cont.)

- **Recurrent Neural Networks (RNNs)** - language modeling, speech recognition/generation
- **Encoder Decoder Architectures** - semantic segmentation, machine translation



Encoder Decoder

- Fully Convolutional Networks for Semantic Segmentation

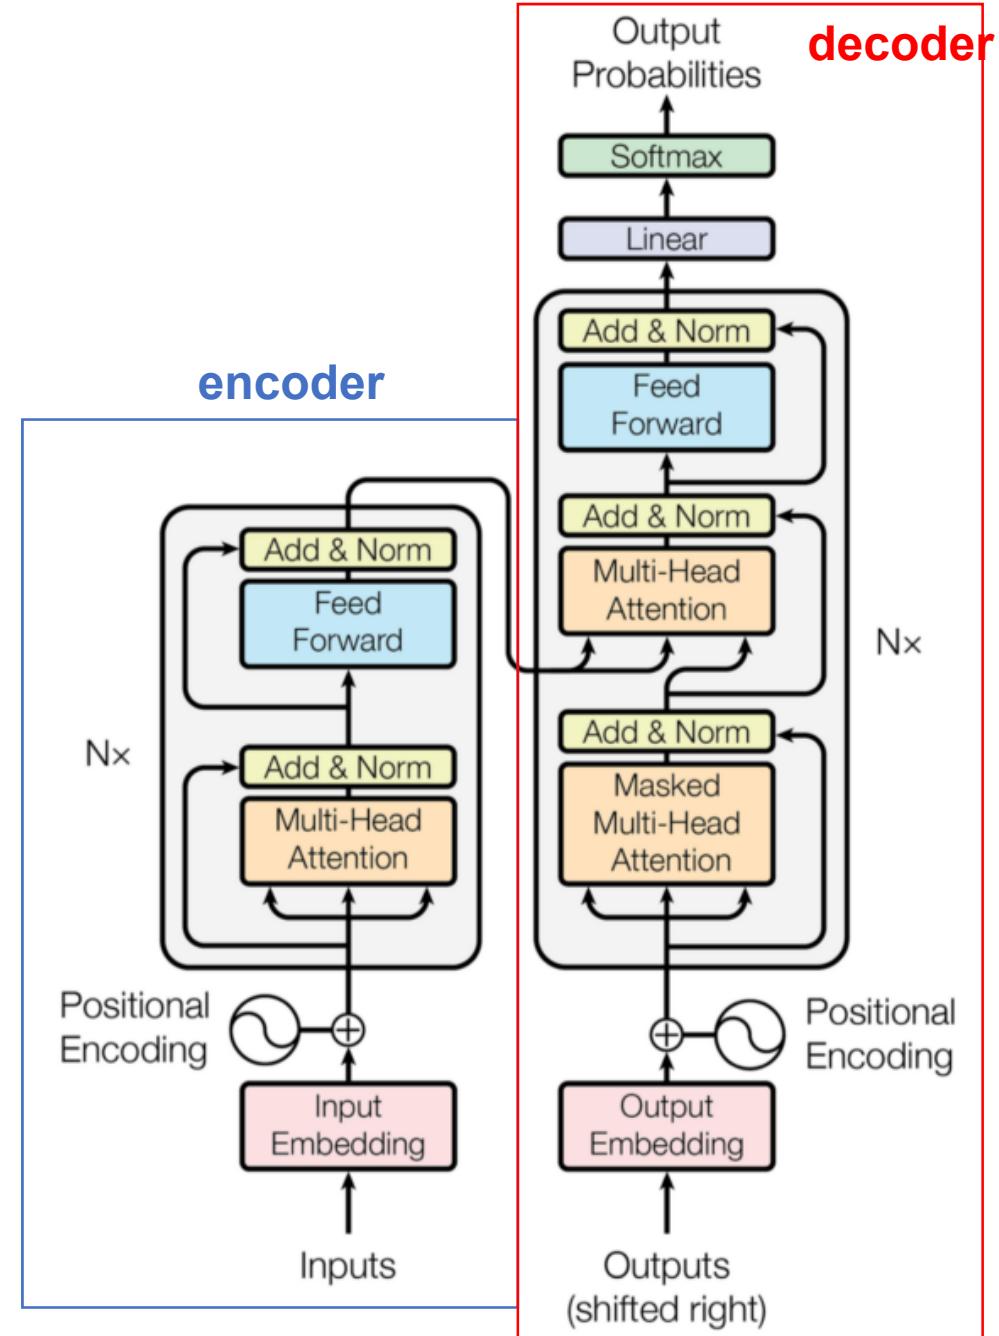


Transformer

$$\begin{array}{c}
 \textbf{x} \quad \textbf{W}^Q \quad \textbf{Q} \\
 \text{---} \times \text{---} = \text{---} \\
 \textbf{x} \quad \textbf{W}^K \quad \textbf{K} \\
 \text{---} \times \text{---} = \text{---} \\
 \textbf{x} \quad \textbf{W}^V \quad \textbf{V} \\
 \text{---} \times \text{---} = \text{---} \\
 \text{softmax} \left(\frac{\textbf{Q} \times \textbf{K}^T}{\sqrt{d_k}} \right) \quad \textbf{V} \quad \textbf{z}
 \end{array}$$

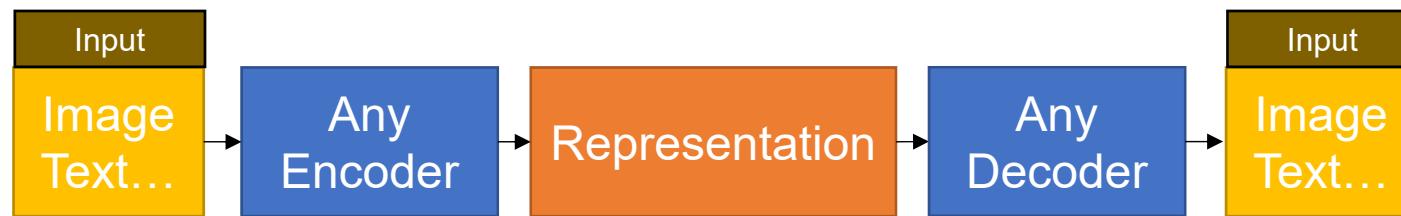
Query Key Value

“Attention Is All You Need”



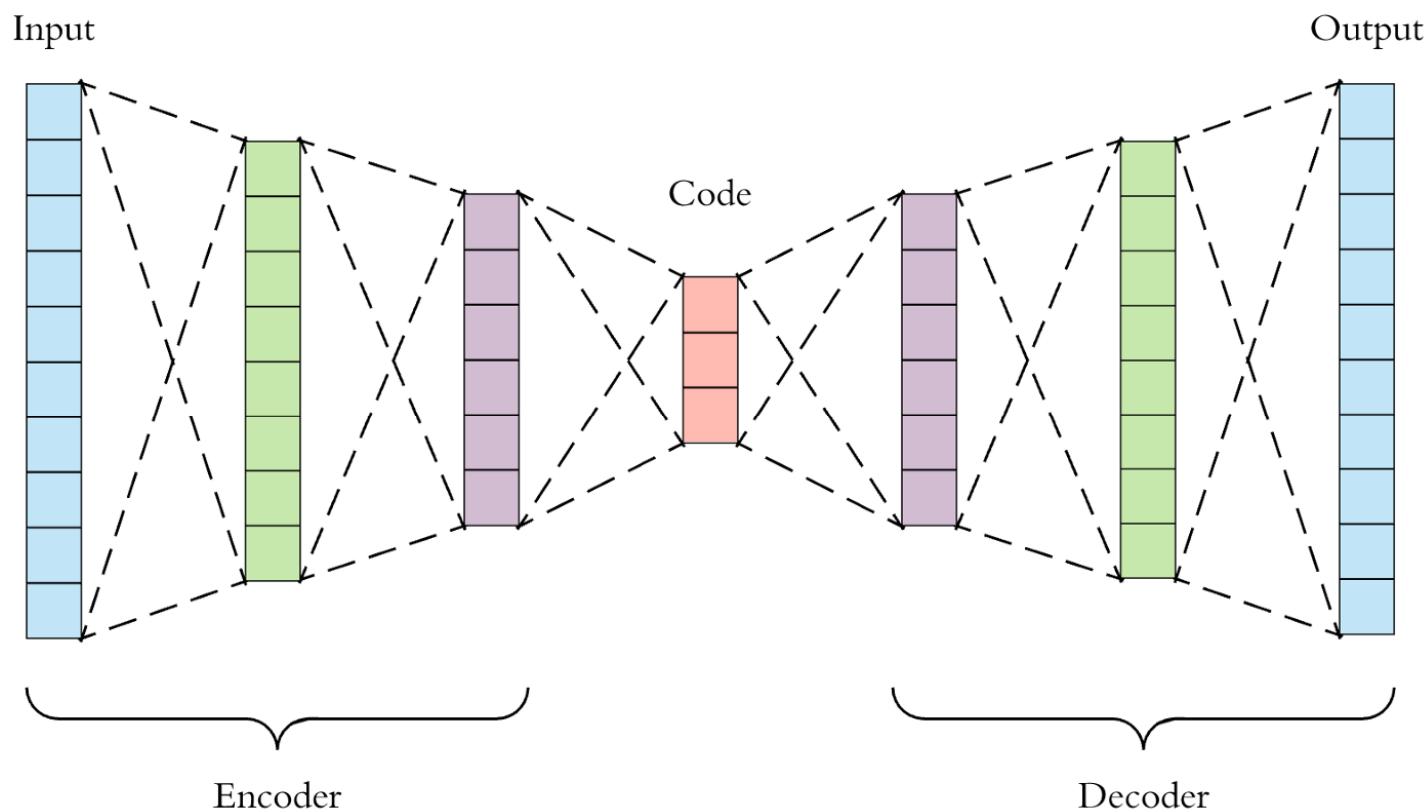
Unsupervised learning

- **Autoencoder** - unsupervised embeddings, denoising

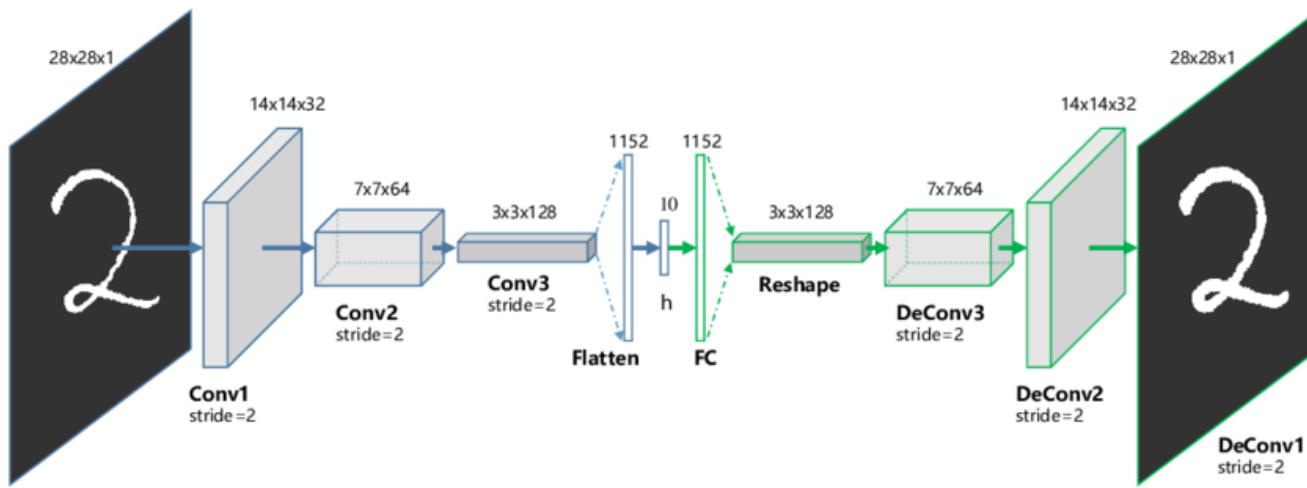


- **Generative Adversarial Networks (GANs)** - generation of realistic images

AutoEncoder

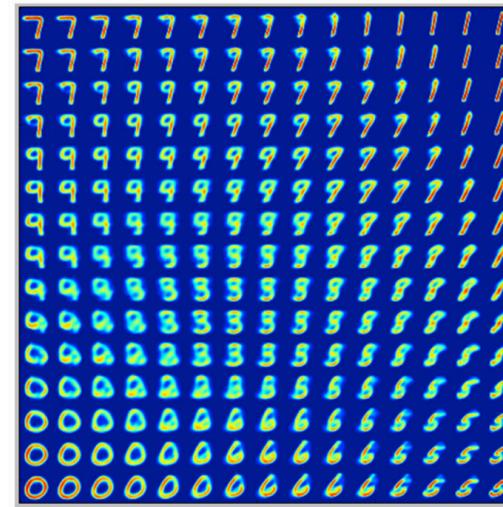
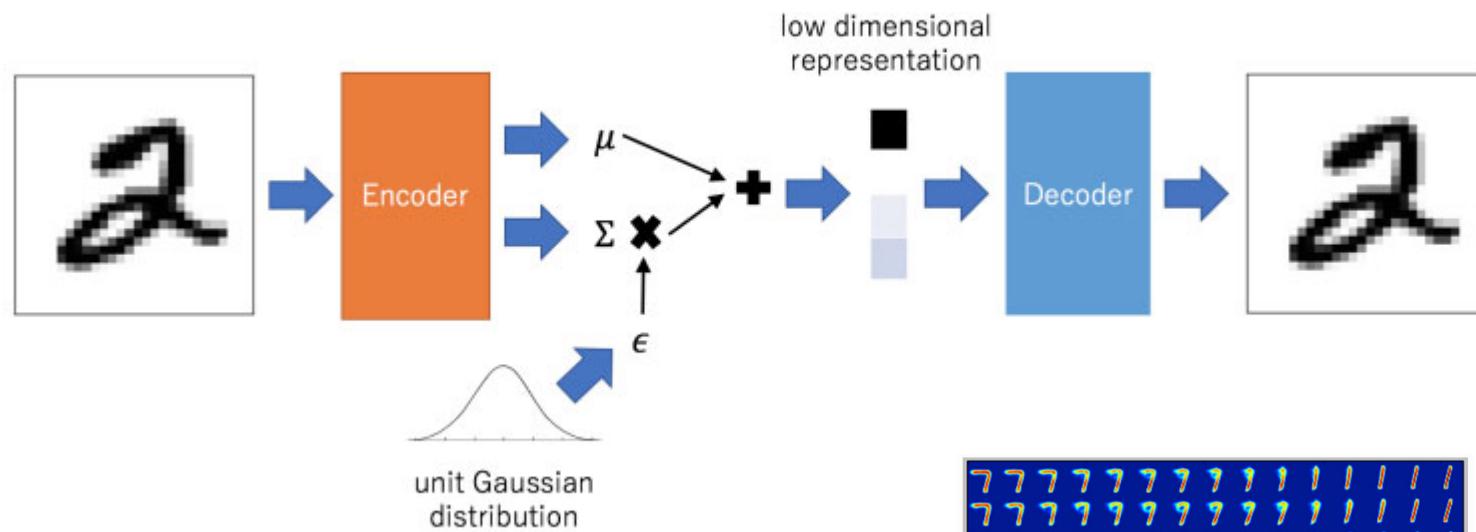


Convolutional AutoEncoder



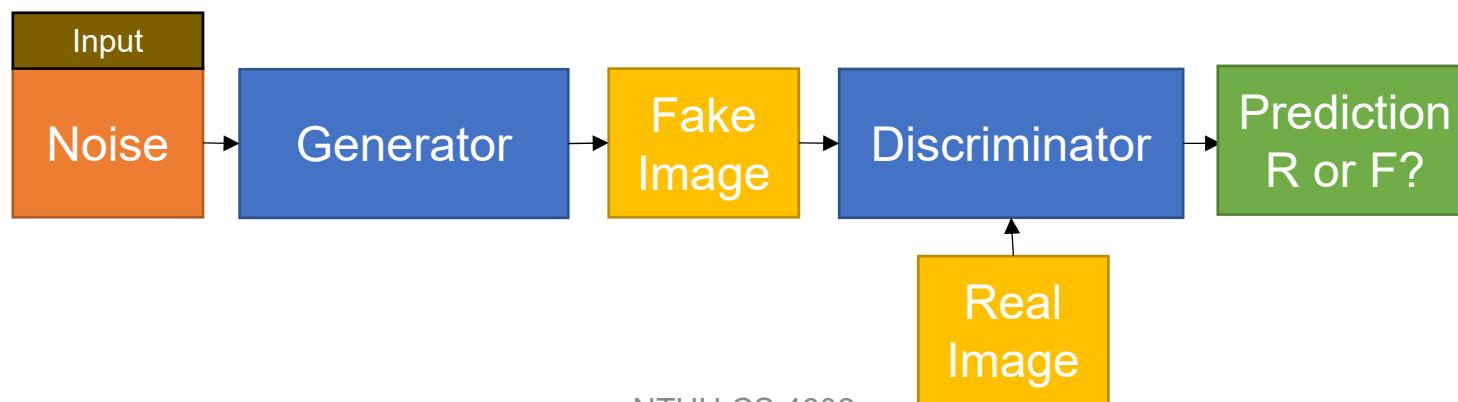
(Guo et al., 2017)

Variational AutoEncoder (VAE)



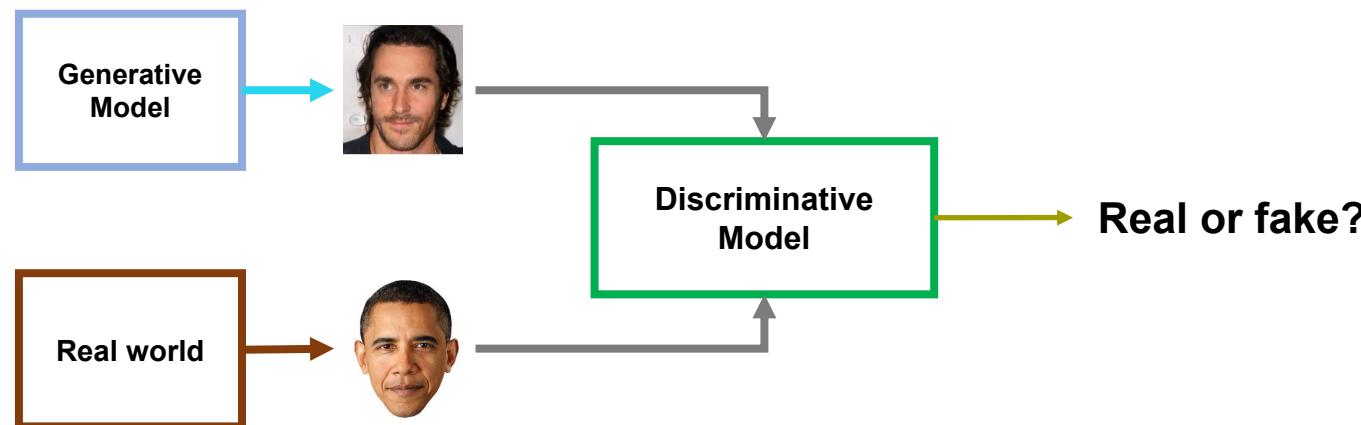
Unsupervised learning

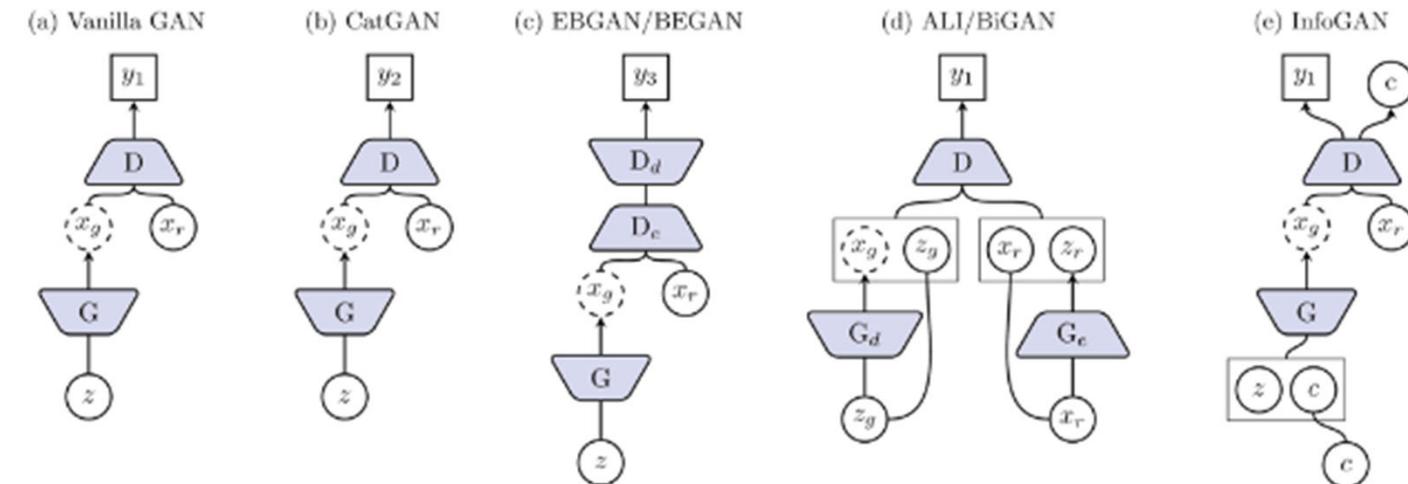
- Autoencoder - unsupervised embeddings, denoising
- **Generative Adversarial Networks (GANs)** - generation of realistic images



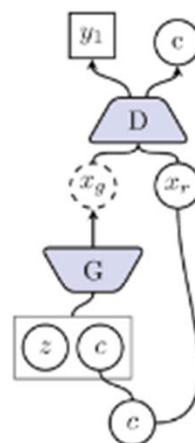
Generative Adversarial Networks (GAN)

- **GAN is the most interesting idea in the last ten years in machine learning.—Yann LeCun, Director, Facebook AI**

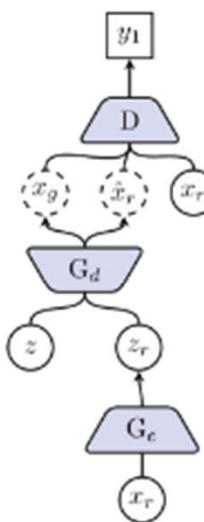




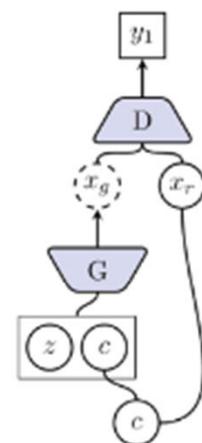
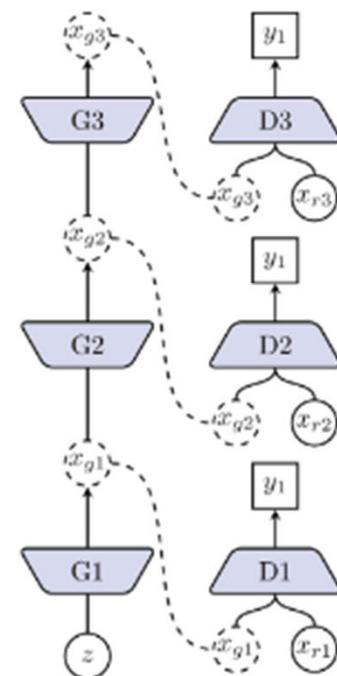
(f) ACGAN



(g) VAEGAN



(h) CGAN

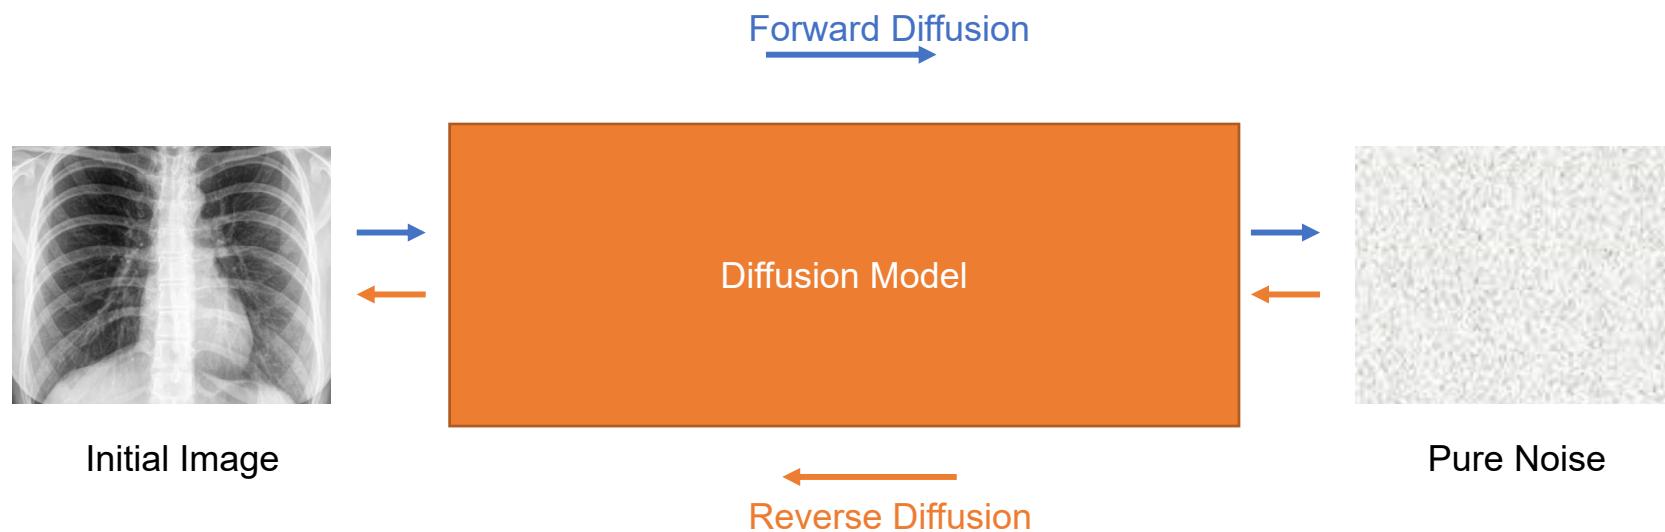
(i) LAPGAN/SGAN
(cascade or stack of GANs)

y_1 real or fake sample

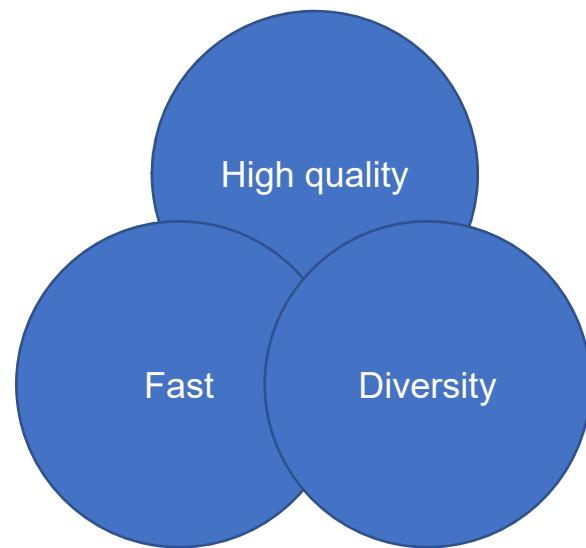
y_2 certain or uncertain class prediction

y_3 real or fake reconstruction loss

Diffusion Models



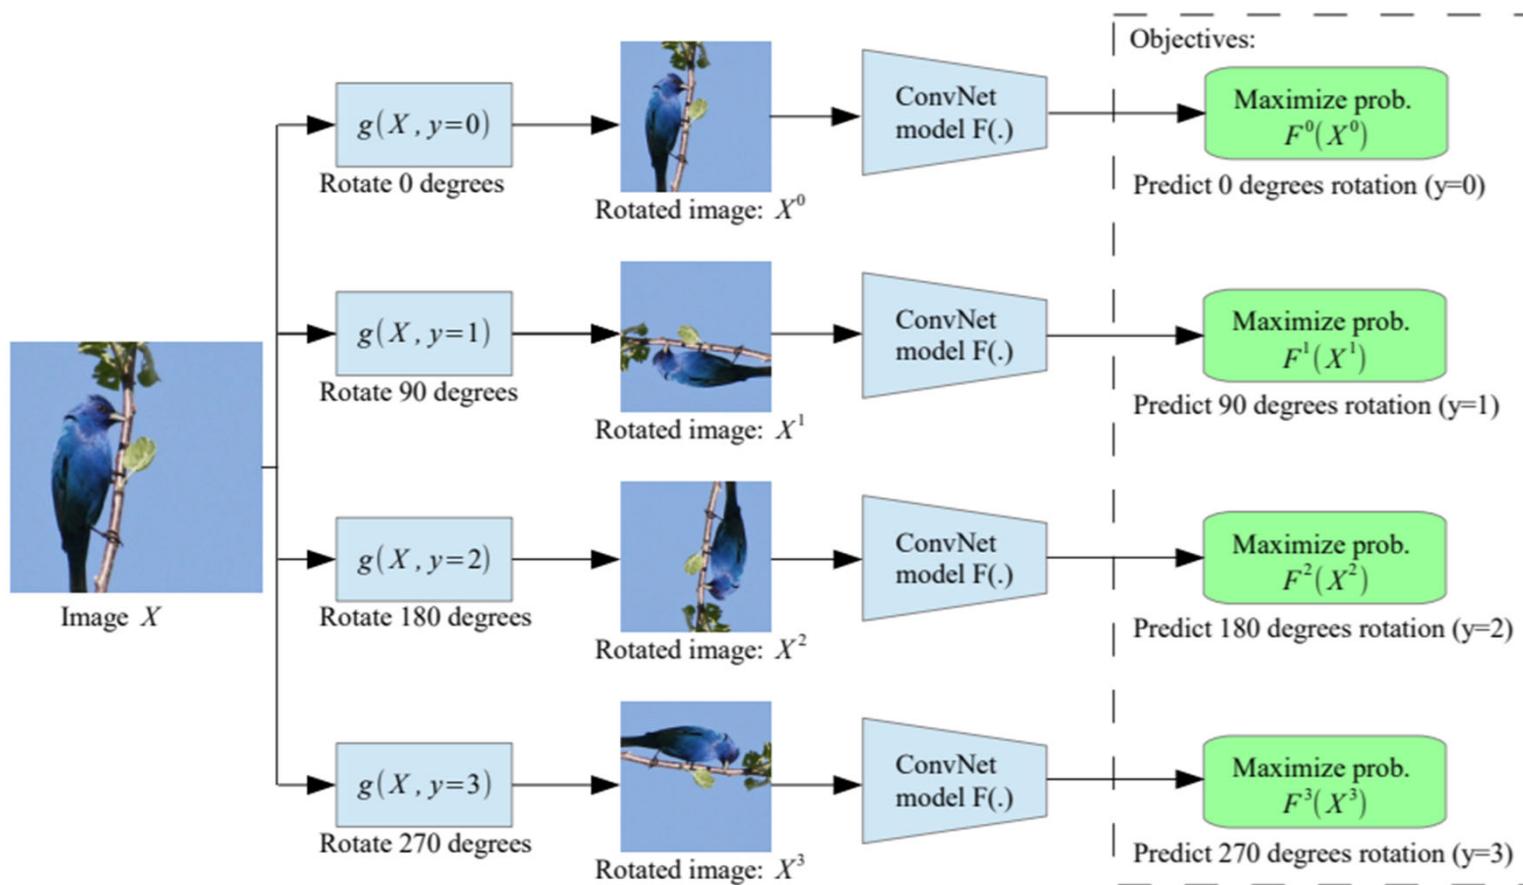
What makes a good generative AI



Self-Supervised Learning (SSL)

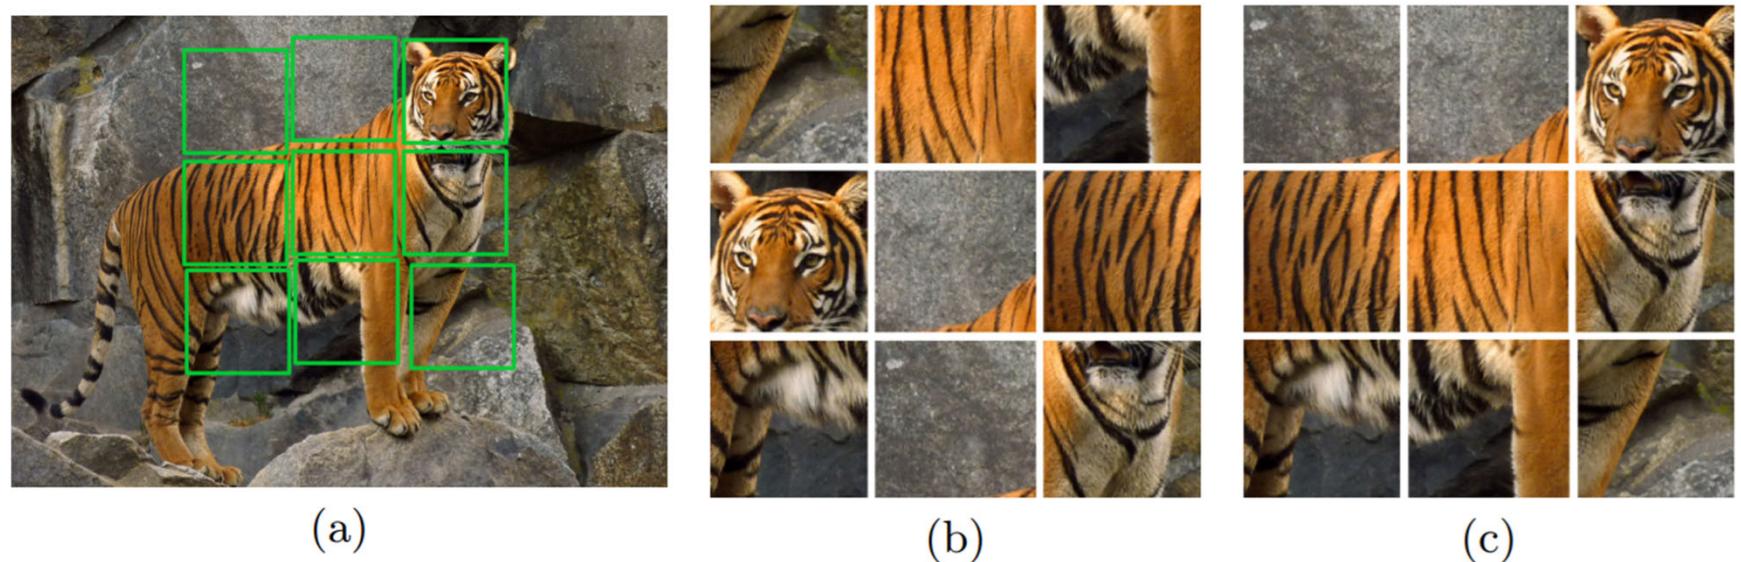
- A representation learning method where a supervised task is created out of the unlabelled data.
- Self-supervised learning is used to reduce the data labelling cost and leverage the unlabelled data pool.
- Some of the popular methods:
 - Rotation
 - Jigsaw Puzzle

Rotation



Unsupervised Representation Learning by Predicting Image Rotations <https://arxiv.org/pdf/1803.07728.pdf>

Jigsaw Puzzle

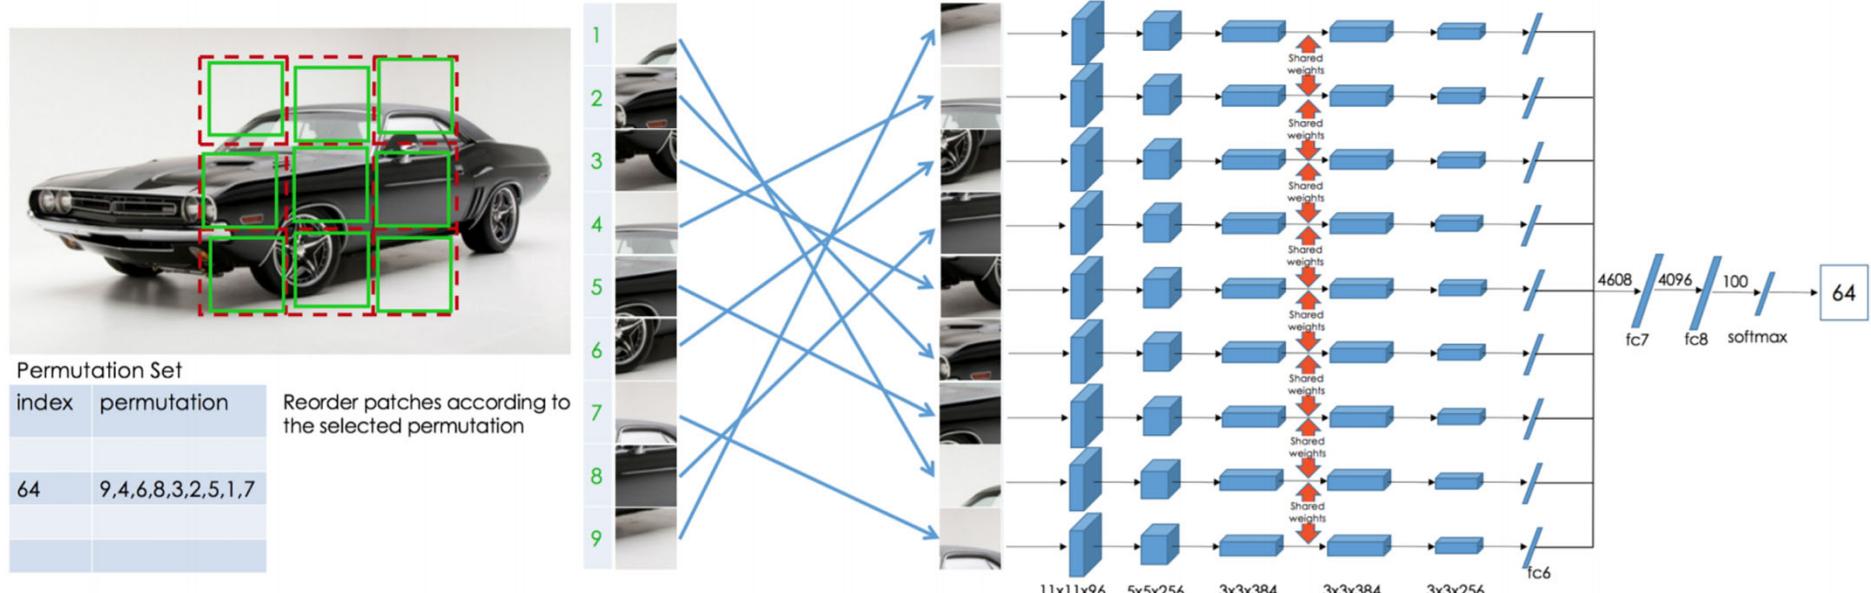


Unsupervised Learning of Visual Representations by Solving Jigsaw
Puzzles <https://arxiv.org/abs/1603.09246>

NTHU CS 4602

30

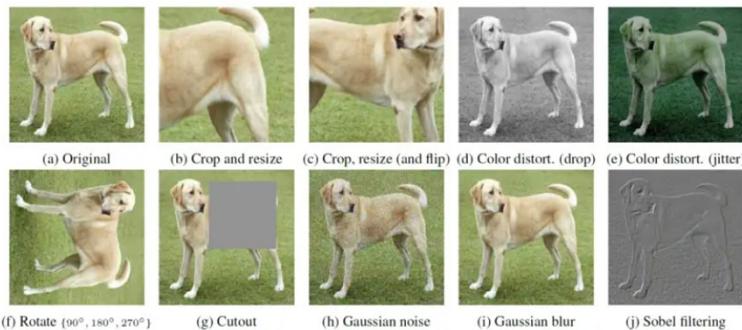
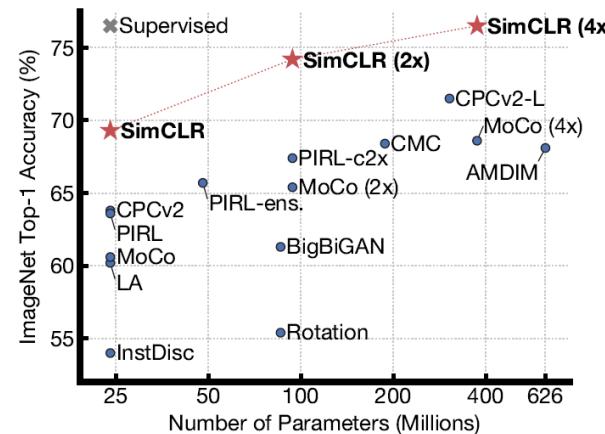
Learning the right positions



Given 9 tiles, there are $9! = 362,880$ possible permutations.

SSL + Contrastive Learning

- SimCLR: A Simple Framework for Contrastive Learning of Visual Representations



GIF for SimCLR (from [Google AI Blog](#))

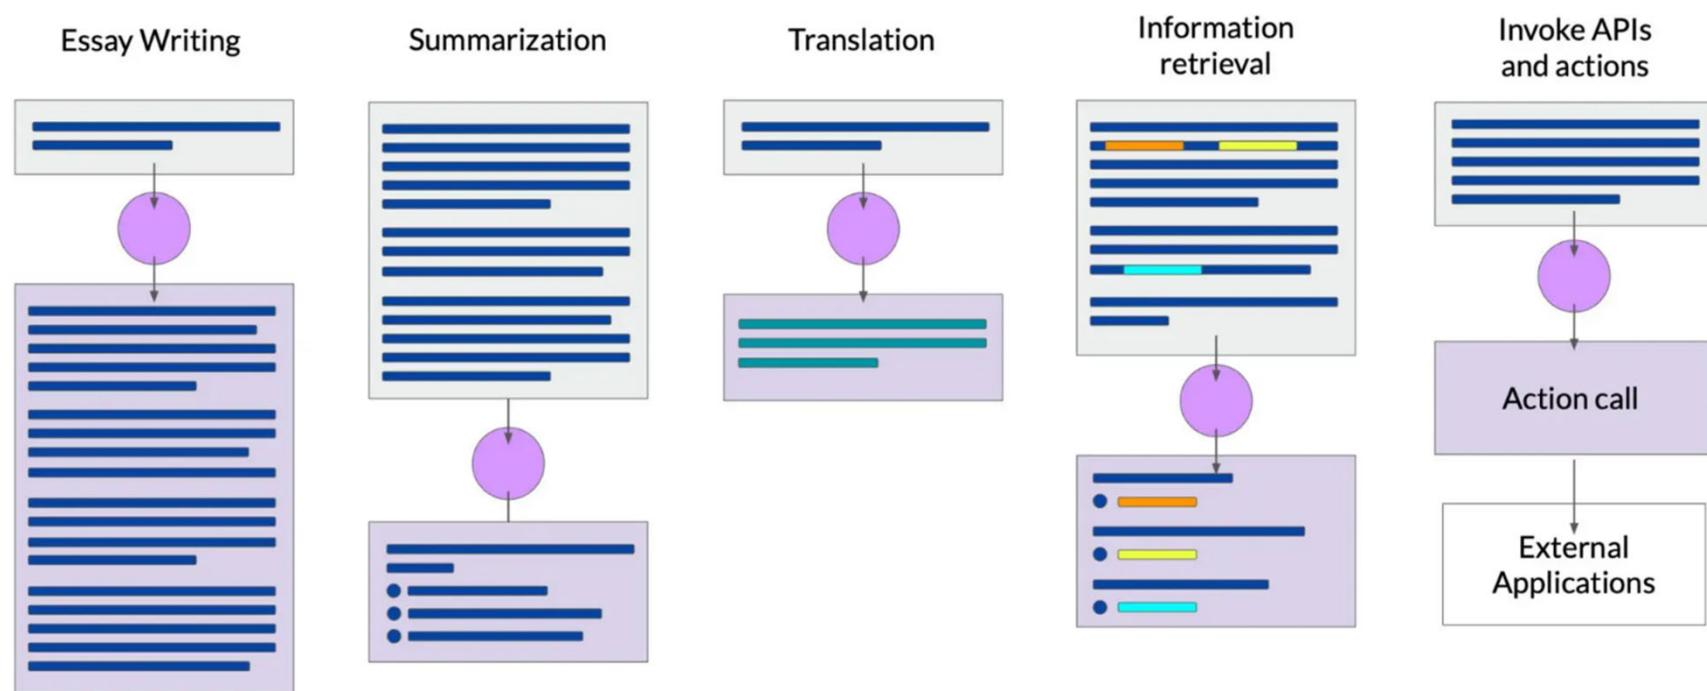
What about LLM (Large Language Model)?

NLP (Natural Language Processing) + Transformer + Self-supervised Learning

Language Model	Description	Is it "Large"?	Emergence
Bag-of-Words Model	Represents text as a set of unordered words, without considering sequence or context.	No	1950s-1960s
N-gram Model	Considers groups of N consecutive words to capture sequence.	No	1950s-1960s
Hidden Markov Models (HMMs)	Represents language as a sequence of hidden states and observable outputs.	No	1980s-1990s
Recurrent Neural Networks (RNNs)	Processes sequential data by maintaining an internal state, capturing context of previous inputs.	No	1990s-2010s
Long Short-Term Memory (LSTM) Networks	Extension of RNNs that captures longer-term dependencies in sequences.	No	2010s
Transformers	Neural network architecture that processes sequences of variable length using a self-attention mechanism.	Yes	2017-Present

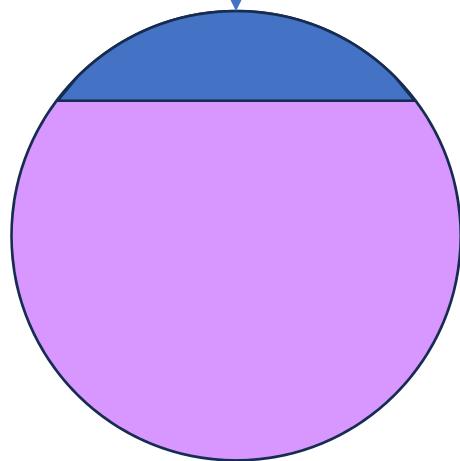


What problem to solve



Source: DeepLearning.AI

Word



Word Embeddings

- Represent words with vectors
- Words with similar meaning tend to occur in similar contexts:
 - The **cat** meowed at me for **food**.
 - The **kitten** meowed at me for **treats**.
- If we use vectors to encode tokens we can attempt to store this meaning.
 - Vectors are the basic inputs for many ML methods.
 - Tokens that are similar in meaning can be positioned as neighbors in the vector space using the right mapping functions.

How to convert words into vectors?

- Initial idea: Let's count the frequency of the words!

<u>Document</u>	<u>the</u>	<u>cat</u>	<u>sat</u>	<u>in</u>	<u>hat</u>	<u>with</u>
the cat sat	1	1	1	0	0	0
the cat sat in the hat	2	1	1	1	1	0
the cat with the hat	2	1	0	0	1	1

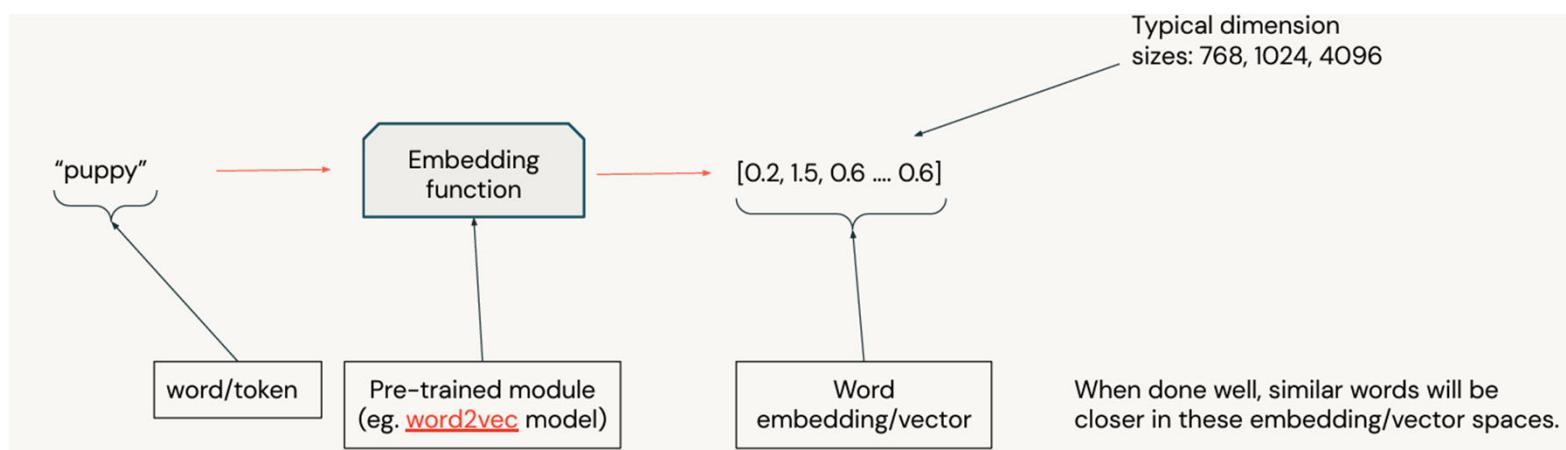
We now have length-6 vectors for each document:

- 'the cat sat' → [1 1 1 0 0 0]
- 'the cat sat in the hat' → [2 1 1 1 1 0]
- 'the cat with the hat' → [2 1 0 0 1 1]

SPARSITY!

Creating dense vector representation

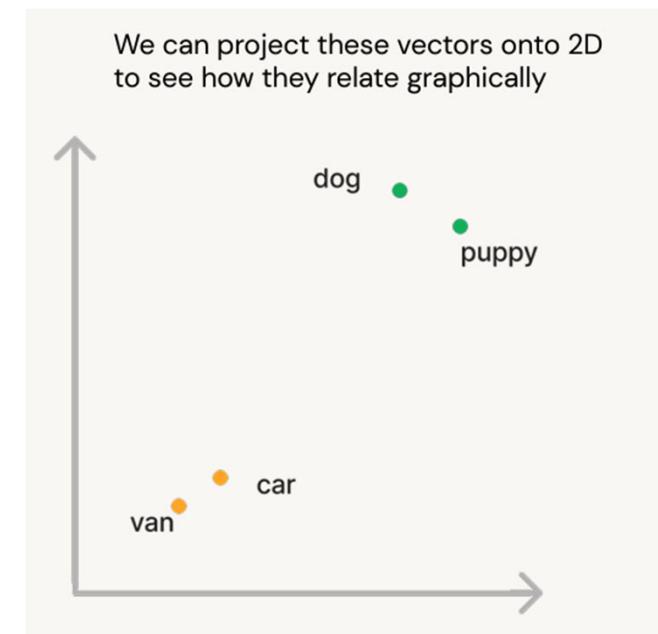
- New idea: Let's give each word a vector representation and use data to build our embedding space.

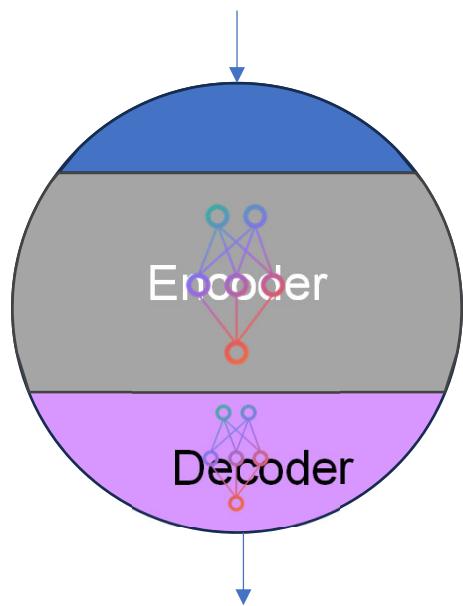


Dense vector representations

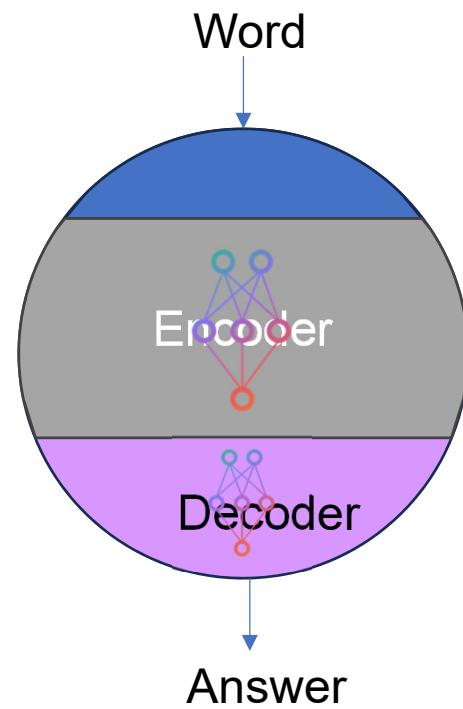
- Visualizing common words using word vectors.

	living being	home	transport	age	
word	N-dimensional word vectors/embeddings				
dog	0.6	0.1	-0.4	0.8
puppy	0.2	1.5	0.6	0.6
car	-0.1	-2.6	0.3	2.4
van	0.9	0.1	-2.5	-1.3

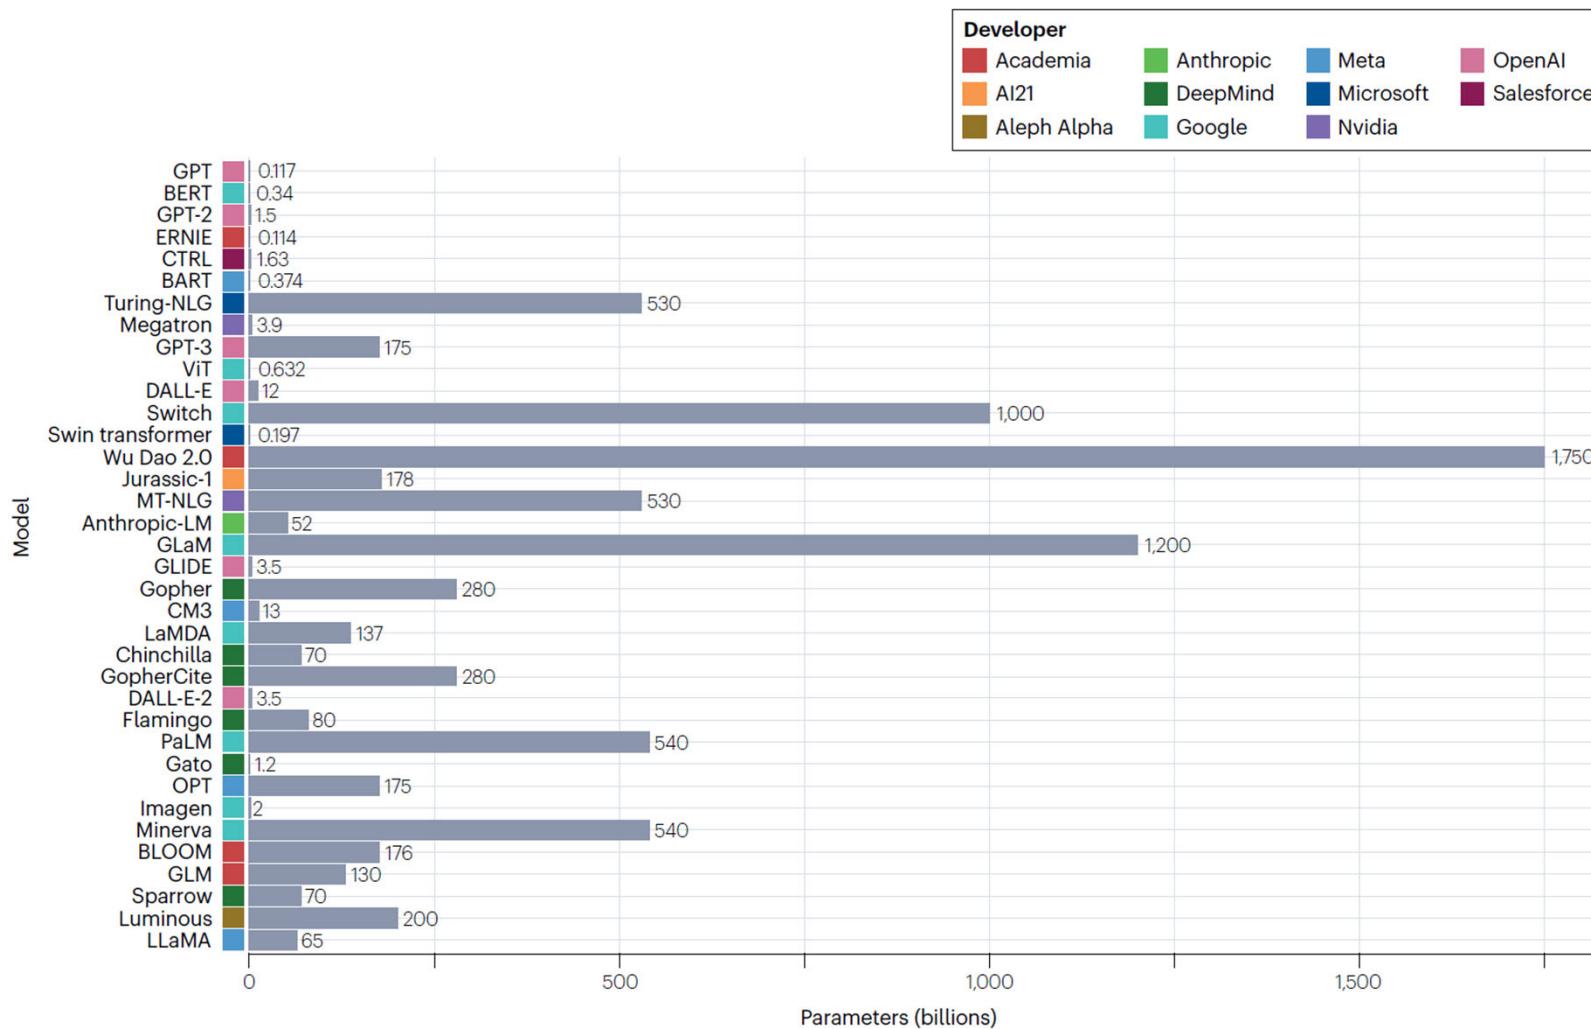


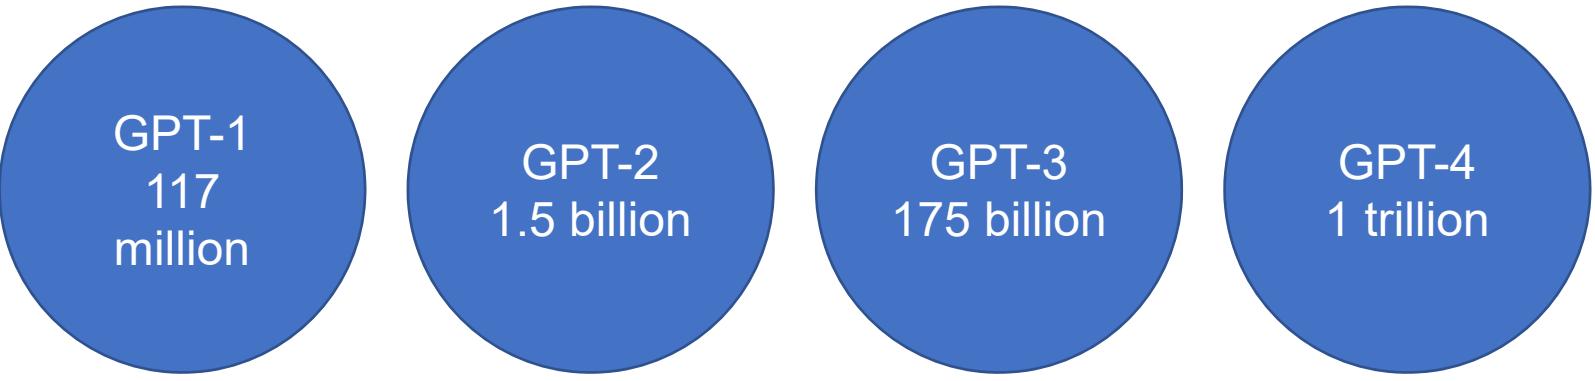


How to train your model?



Parameters of LLM





GPT-1
117
million

GPT-2
1.5 billion

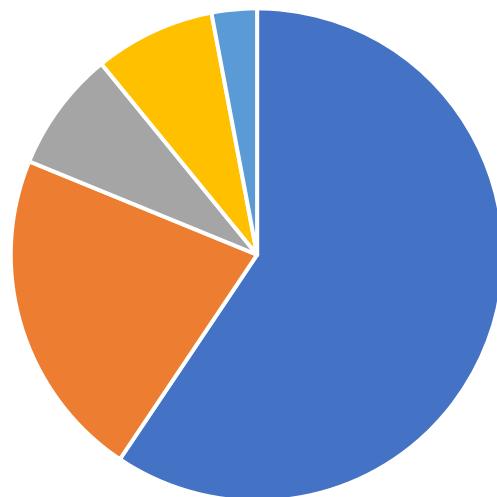
GPT-3
175 billion

GPT-4
1 trillion

We need data!

GPT-3

Data source

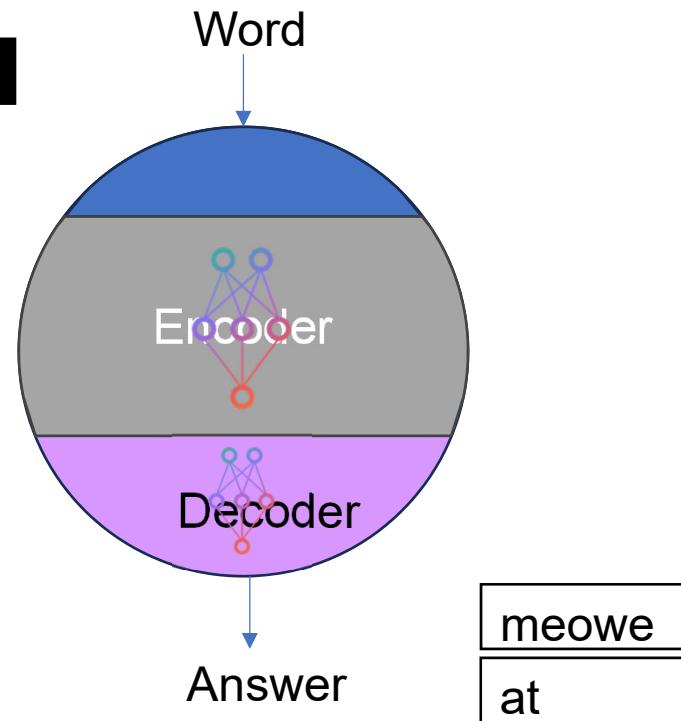


- Common Crawl
- WebText2
- Book1
- Book2
- Wikipedia

How to train your model?

The kitten █ at
me for treats.

The kitten meowed █
me for treats.

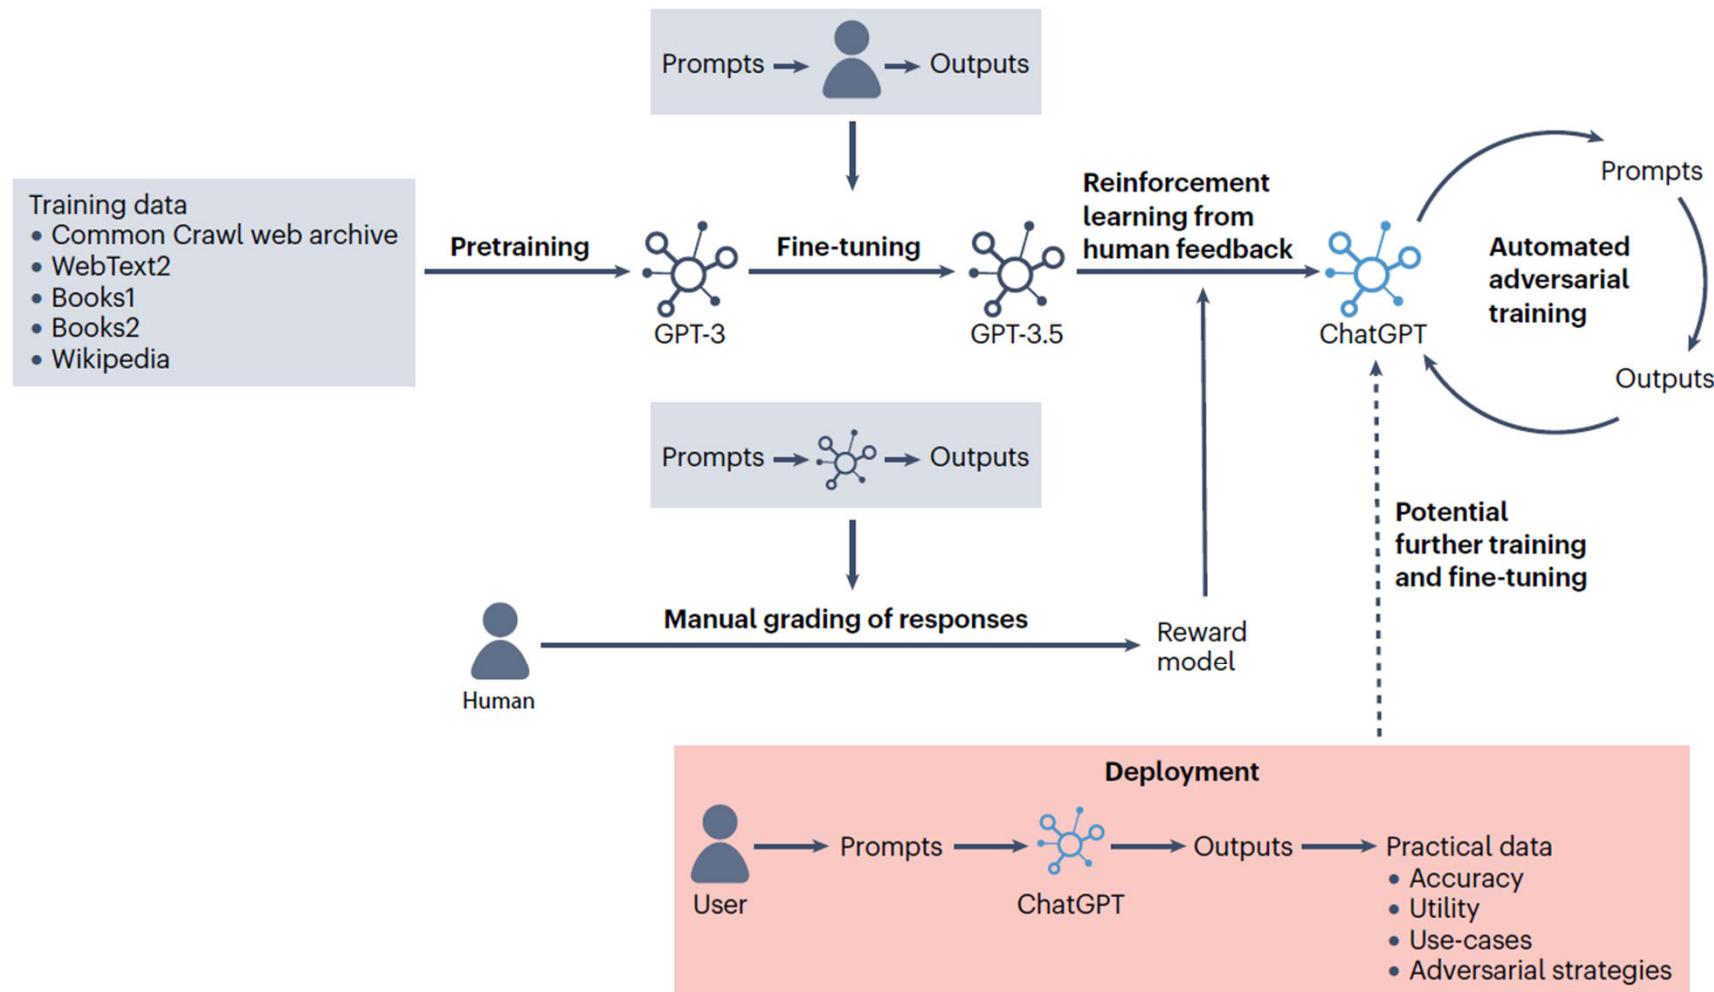


Training schemata

- Masked language modeling
 - cloze tasks: predicting masked tokens in a sequence
- Permuted language modeling
 - language modeling with randomly sampled input tokens
- Denoising autoencoding
 - Recovering undistorted inputs after intentional corruption
- Next-sentence prediction
 - Distinguishing whether sentences are contiguous or not.
- Examples
 - Gato, DALL-E, Enhanced Language Representations with Informative Entities (ERNIE), Bidirectional Encoder Representations from Transformers (BERT) and Bidirectional and Auto-regressive Transformers (BART)

From LLM to generative AI chatbot

- Fine-tuning involved exposing GPT-3 to prompts and responses produced by human researchers acting the part of an application user and AI assistant
- **Reinforcement learning from human feedback' (RLHF)** was conducted using a reward model trained on data generated by human graders tasked with ranking GPT-3.5 responses to a set of queries.



Reinforcement Learning

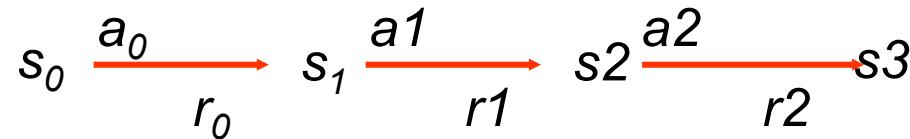
- **Deep Reinforcement Learning** - game playing, robotics in simulation, self-play, neural architecture search

Reinforcement Learning

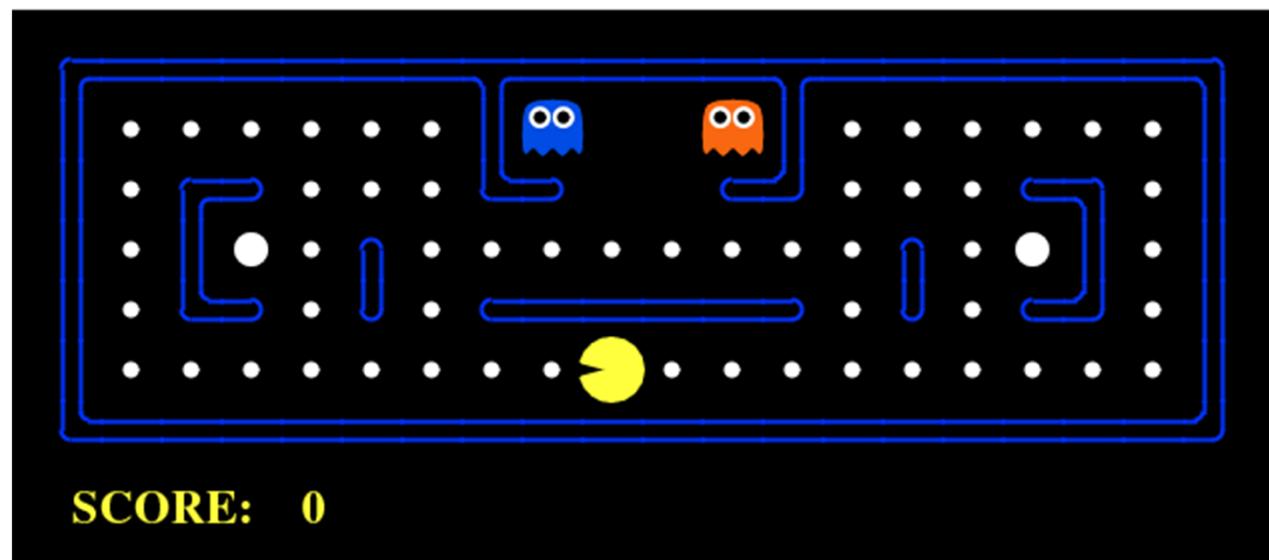
- Supervised Learning: Immediate feedback (labels provided for every input).
- Unsupervised Learning: No feedback (no labels provided).
- Reinforcement Learning: Delayed scalar feedback (a number called reward).
- RL deals with agents that must sense & act upon their environment.
- Examples:
 - A robot cleaning the room and recharging its battery
 - Robot-soccer
 - Modeling the economy through rational agents
 - Learning how to fly a helicopter
 - Scheduling planes to their destinations

The Big Picture

- S – set of states
- A – set of actions
- $T(s,a,s') = P(s'|s,a)$ – the probability of transition from s to s' given action a
- $R(s,a)$ – the expected reward for taking action a in state s



Your action influences the state of the world which determines its reward



The Task

- To learn an optimal *policy* that maps states of the world to actions of the agent.
I.e., if this patch of room is dirty, I clean it. If my battery is empty, I recharge it.

$$\pi : S \rightarrow A$$

- What is it that the agent tries to optimize?
 - the total discounted future reward:

$$\begin{aligned} V^\pi(s_t) &= r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \\ &= \sum_{i=0}^{\infty} \gamma^i r_{t+i}, \quad 0 \leq \gamma < 1 \end{aligned}$$

Note: immediate reward is worth more than future reward.

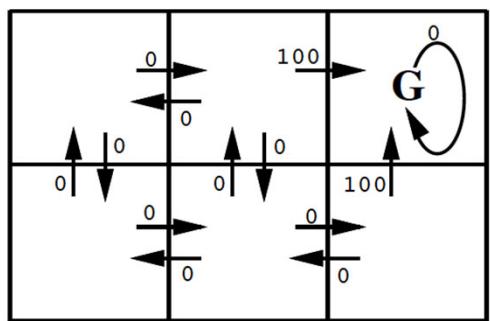
Value Function

- Let's say we have access to the optimal value function that computes the total future reward $V^*(s)$
- The optimal policy $\pi^*(s)$ is chosen by maximizing:

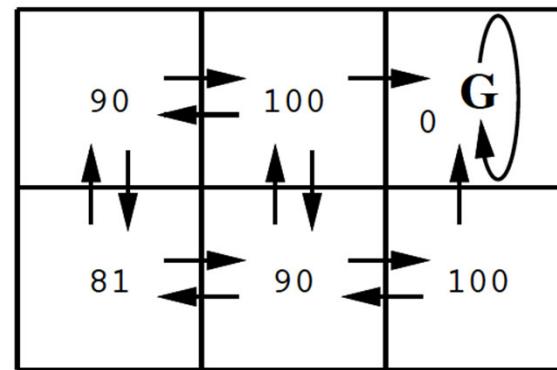
$$\pi^*(s) = \underset{a}{\operatorname{argmax}} \left[r(s, a) + \gamma V^*(\delta(s, a)) \right]$$

- We assume that we know what the reward will be if we perform action "a" in state "s": $r(s, a)$
- We also assume we know what the next state of the world will be if we perform action "a" in state "s":
 $s_{t+1} = \delta(s_t, a)$

Example: Find your way to G

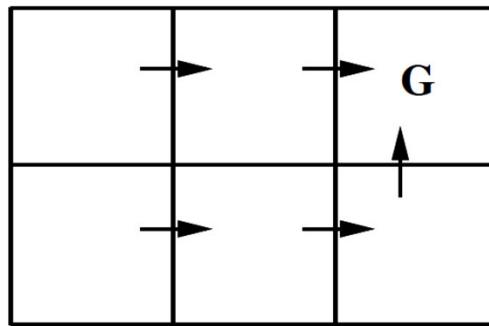


$r(s, a)$ (immediate reward) values



$V^*(s)$ values

$$\gamma = 0.9$$



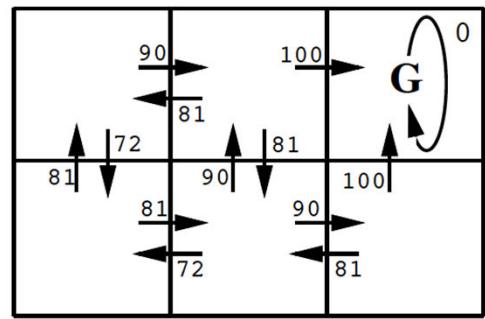
One optimal policy

Q-Function

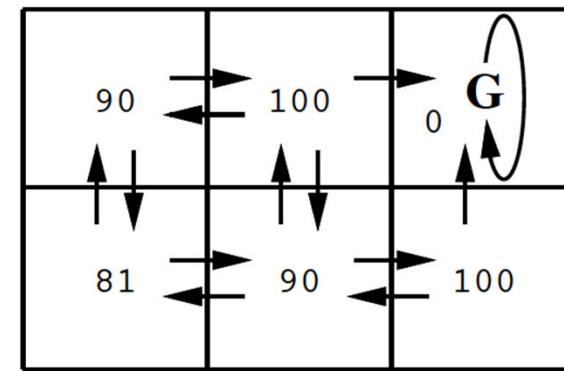
- One approach to RL to estimate $V^*(s)$.
$$V^*(s) \leftarrow \max_a [r(s,a) + \gamma V^*(\delta(s,a))]$$
- However, this approach requires you to know $r(s,a)$ and $\delta(s,a)$.
- We need a function that directly learns good state-action pairs, i.e. what action should I take in this state. We call this $Q(s,a)$.
- Given $Q(s,a)$ it is now trivial to execute the optimal policy, *without knowing $r(s,a)$ and $\delta(s,a)$* . We have:

$$\pi^*(s) = \arg \max_a Q(s,a)$$

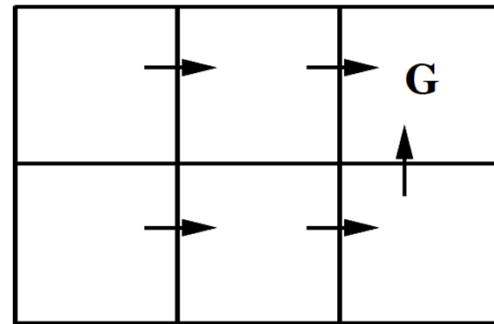
$$V^*(s) = \max_a Q(s,a)$$



$Q(s, a)$ values



$V^*(s)$ values



One optimal policy

$$\pi^*(s) = \arg \max_a Q(s, a)$$

$$V^*(s) = \max_a Q(s, a)$$

Q-Learning

$$\begin{aligned} Q(s, a) &\equiv r(s, a) + \gamma V^*(\delta(s, a)) \\ &= r(s, a) + \gamma \max_{a'} Q(\delta(s, a), a') \end{aligned}$$

- This still depends on $r(s, a)$ and $\delta(s, a)$.
- However, imagine the robot is exploring its environment, trying new actions as it goes.
- At every step it receives some reward “ r ”, and it observes the environment change into a new state s' for action a .
How can we use these observations, (s, a, s', r) to learn a model?

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a') \quad s' = s_{t+1}$$

- This equation continually estimates Q at state s consistent with an estimate of Q at state s' : **temporal difference (TD-Q) learning**.
- Do an update after each state-action pair.

3 steps in Q-learning

1. Agent starts in a state (s_1) takes an action (a_1) and receives a reward (r_1)
2. Agent selects action by referencing Q-table with highest value (max) OR by random (epsilon, ε)
3. Update q-values

Initialized

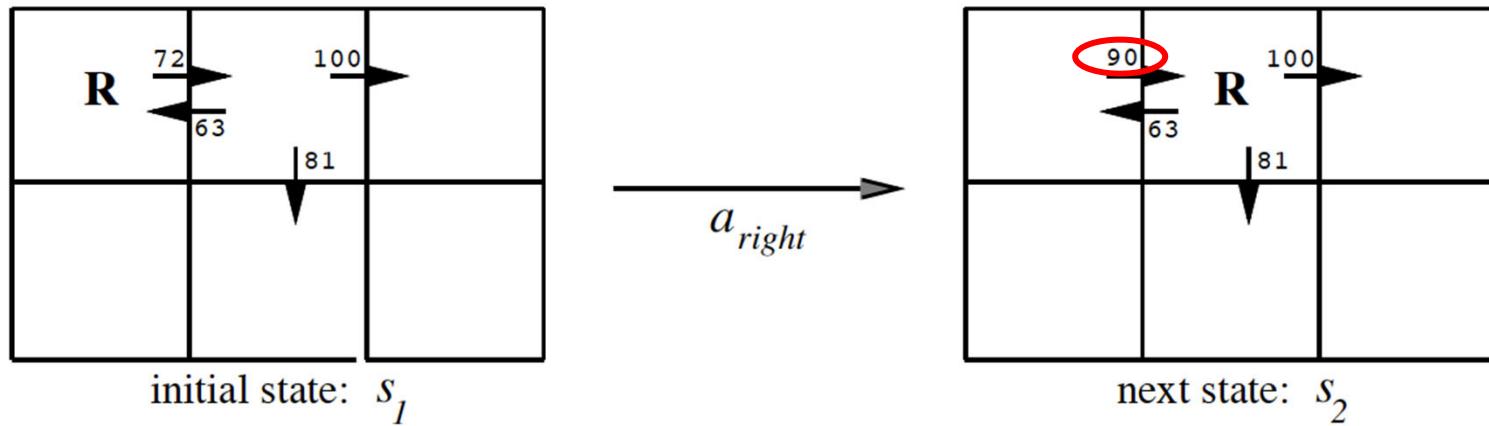
Q-Table		Actions					
		South (0)	North (1)	East (2)	West (3)	Pickup (4)	Dropoff (5)
States	0	0	0	0	0	0	0
.
.
.
States	327	0	0	0	0	0	0
.
.
.
States	499	0	0	0	0	0	0



Q-Table		Actions					
		South (0)	North (1)	East (2)	West (3)	Pickup (4)	Dropoff (5)
States	0	0	0	0	0	0	0
.
.
States	328	-2.30108105	-1.97092096	-2.30357004	-2.20591839	-10.3607344	-8.5583017
.
.
States	499	9.96984239	4.02706992	12.96022777	29	3.32877873	3.38230603

<https://en.wikipedia.org/wiki/Q-learning>

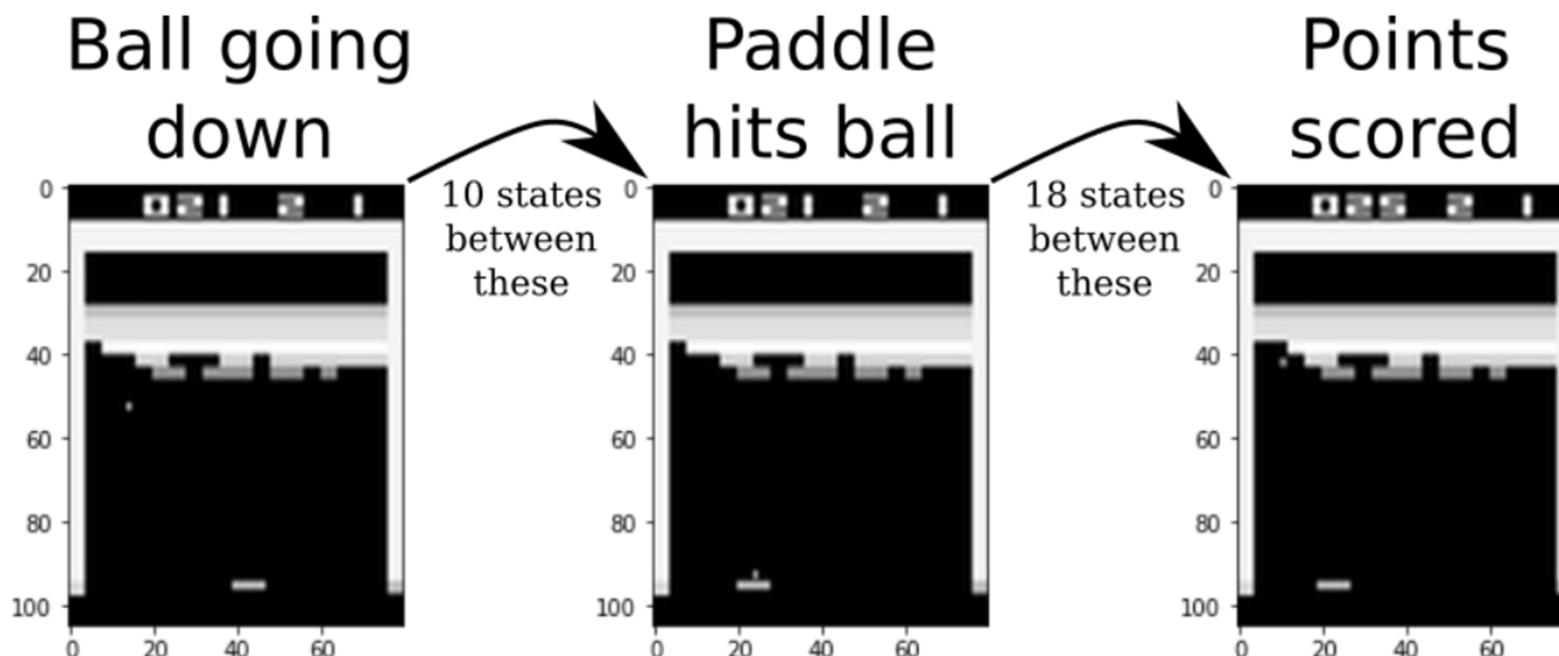
Example Q-Learning



$$\begin{aligned}\hat{Q}(S_1, a_{right}) &\leftarrow r + \gamma \max_{a'} \hat{Q}(S_2, a') \\ &\leftarrow 0 + 0.9 \max(63, 81, 100) \\ &\leftarrow 90\end{aligned}$$

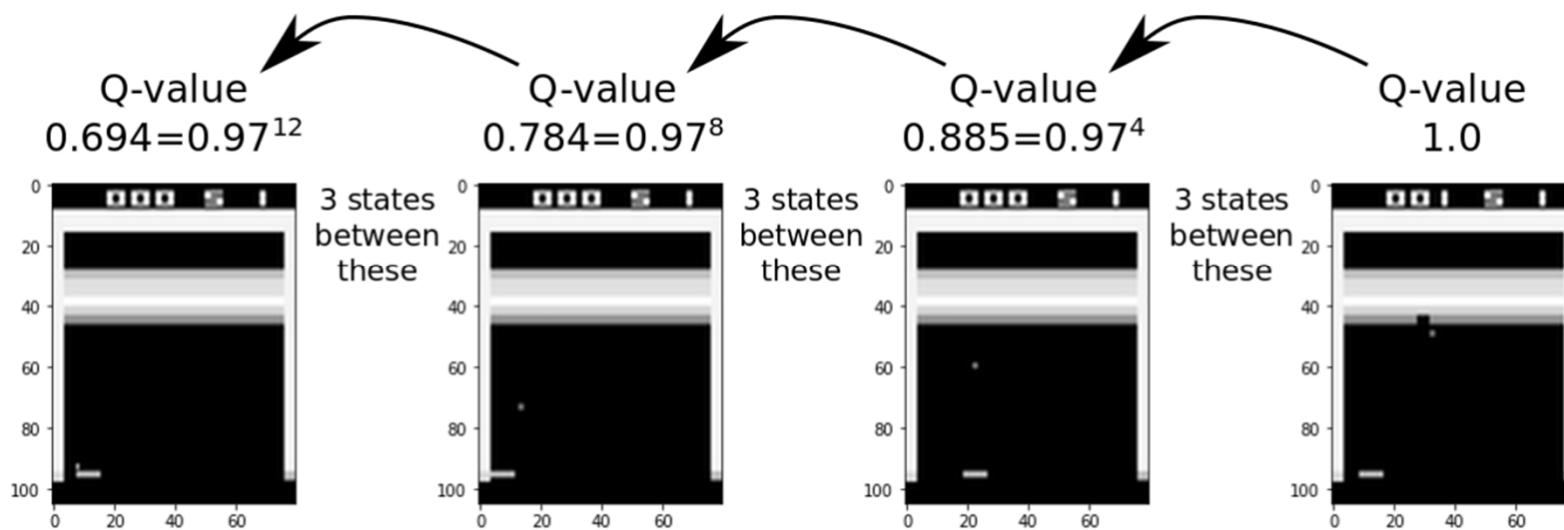
Q-learning propagates Q-estimates 1-step backwards

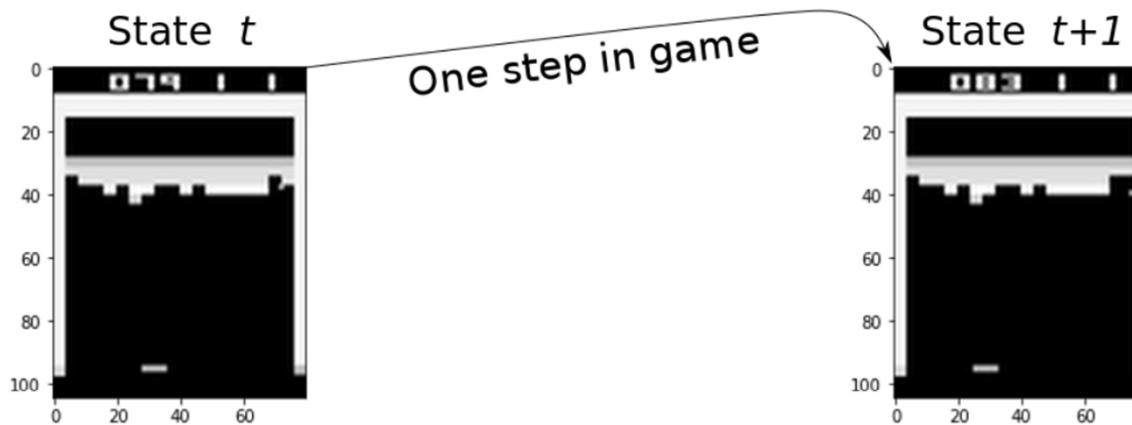
Example: Atari Breakout



https://colab.research.google.com/github/Hvass-Labs/TensorFlow-Tutorials/blob/master/16_Reinforcement_Learning.ipynb#scrollTo=CT6gxbxIJ-V_

$$Q(s_t, a_t) \leftarrow \underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of future rewards}}$$

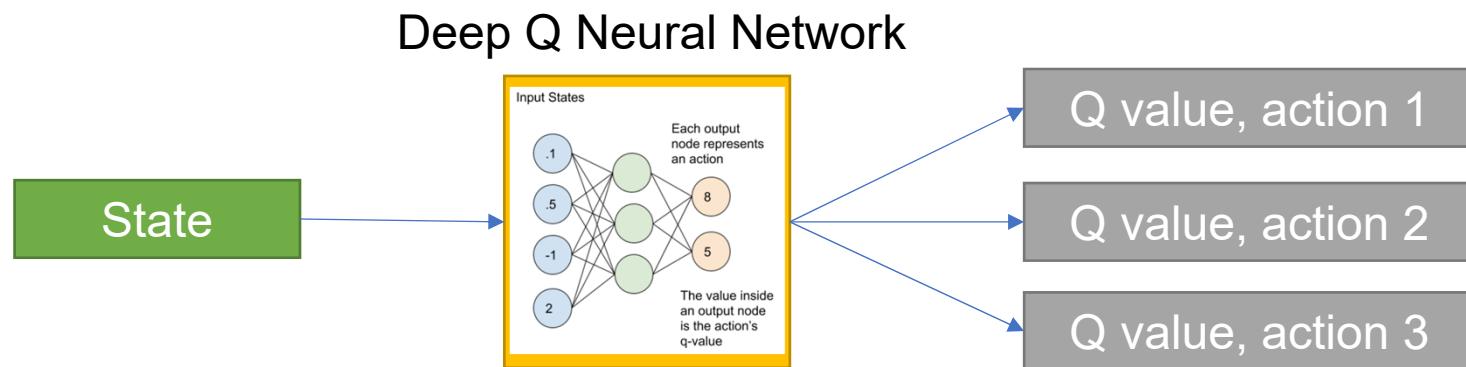
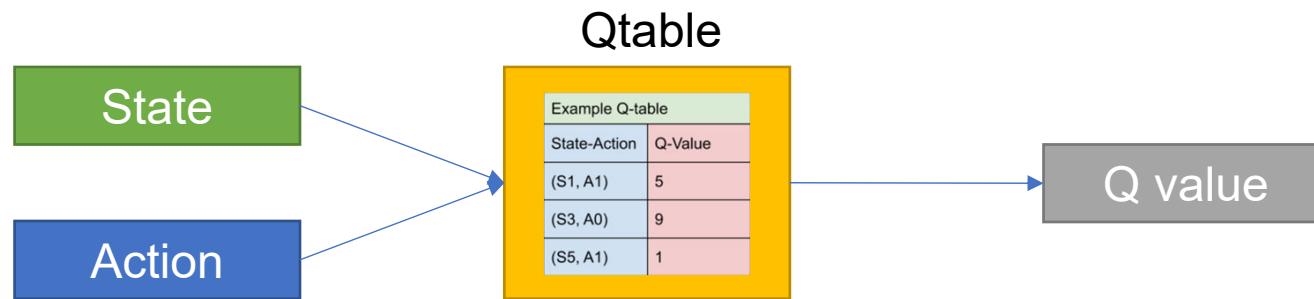




Action:	Q-Values:	Action:	Q-Values:
=====	=====	=====	=====
NOOP	2.900 (Action Taken)	NOOP	1.814
FIRE	2.622	FIRE	1.777
RIGHT	2.078	RIGHT	1.825
LEFT	2.266	LEFT	1.697
RIGHTFIRE	2.183	<i>Update Q-Value</i>	RIGHTFIRE 1.830 (Max Value)
LEFTFIRE	2.780	$1 + 0.97 \times 1.830 = 2.775$	

$$Q(state_t, NOOP) \leftarrow \underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount}} \cdot \underbrace{\max_a Q(state_{t+1}, a)}_{\text{estimate of future rewards}} = 1.0 + 0.97 \cdot 1.830 \simeq 2.775$$

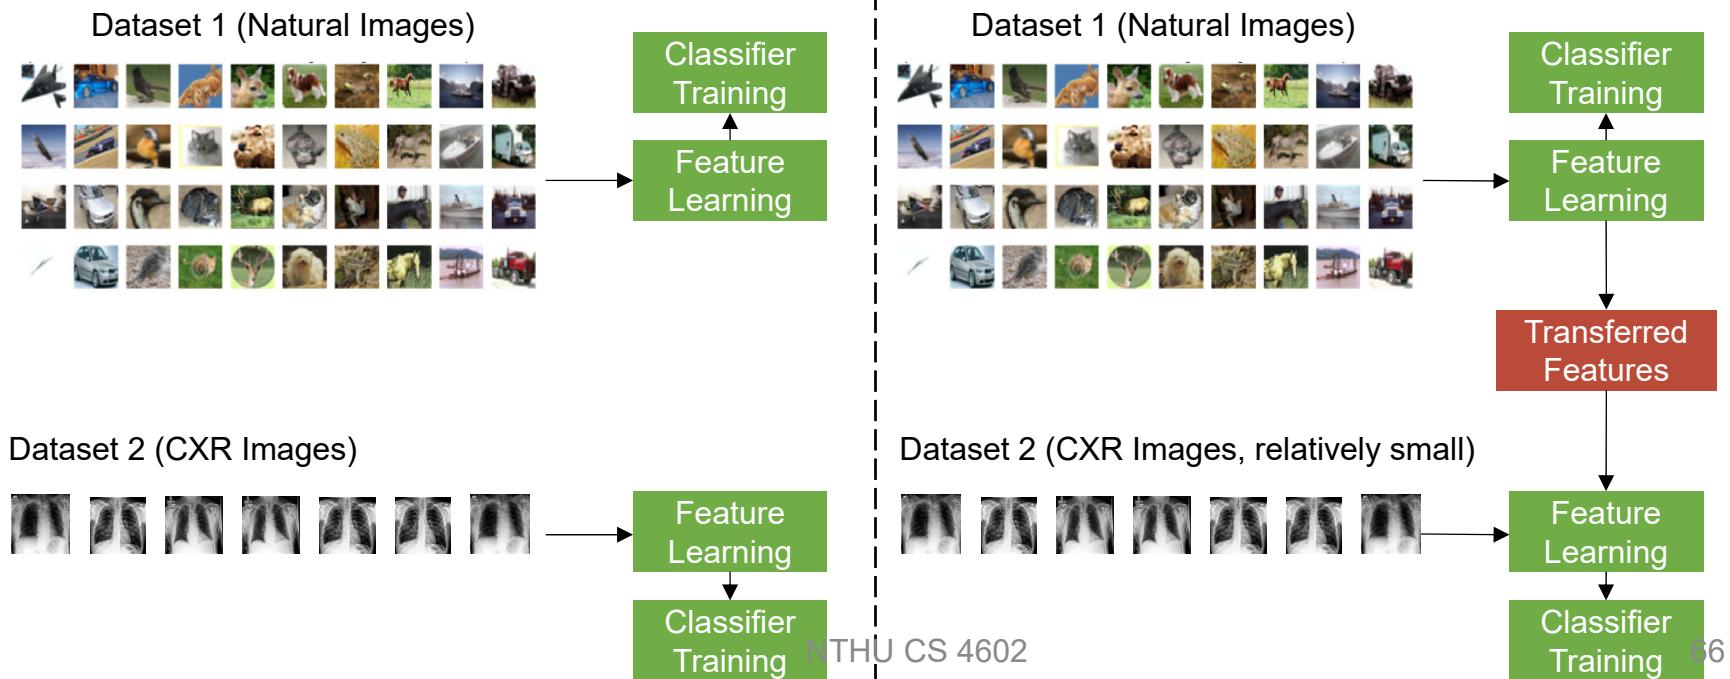
RL + DL = Deep Q-Network (DQN)



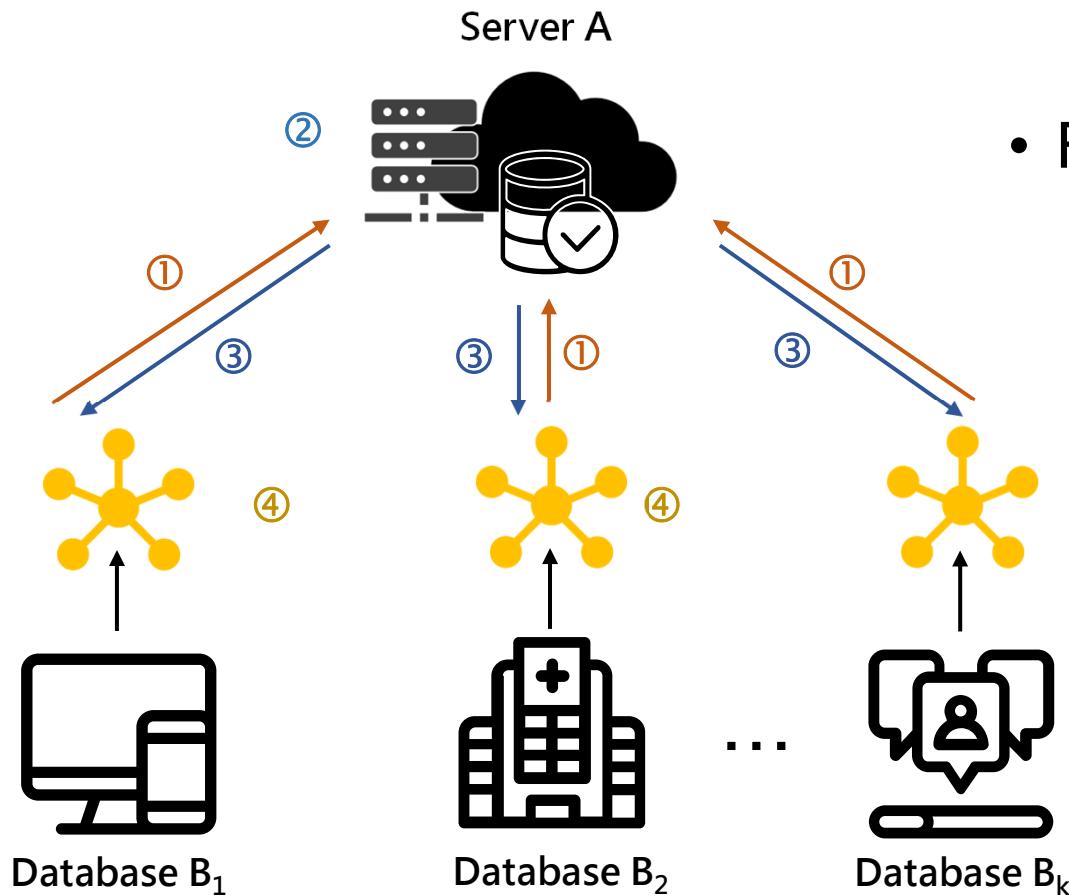
Transfer learning

- Improvement of learning for new task through the transfer of knowledge from similar task that has been trained.

Traditional learning vs. Transfer learning



Federated learning



- Federated Learning Steps

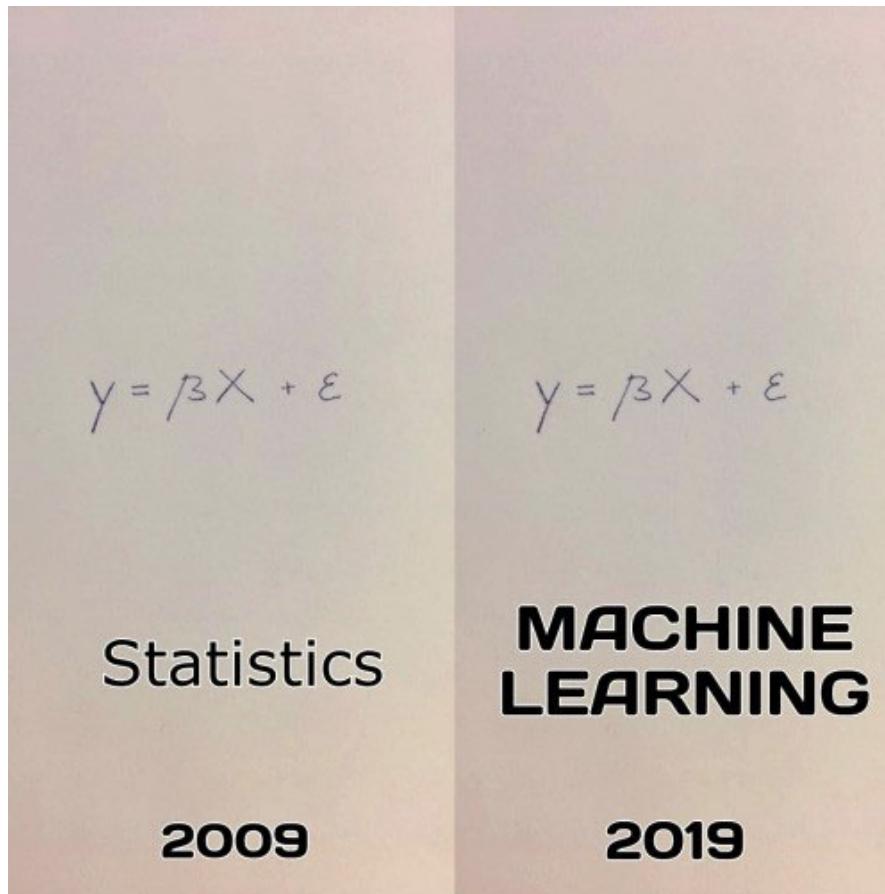
- ① Sending encrypted gradients
- ② Secure aggregation
- ③ Sending back model updates
- ④ Updating models

Konečný, Jakub, et al. "Federated learning: Strategies for improving communication efficiency." *arXiv preprint arXiv:1610.05492* (2016).

Learning, learning, learning

- Supervised learning
 - **FFNNs**
 - **CNNs**
 - **RNNs**
 - **Encoder Decoder**
- Unsupervised learning
 - **Autoencoder**
 - **GAN**
- Self-Supervised learning
- Reinforcement learning
- Transfer learning
- Federated learning

Questions?



#10yearchallenge